

1 GPU-accelerated Bayesian calibration of a variational
2 multiphase model of multispecies biofilms [final title
3 after the results]

4 Keisuke Nishioka^{a,d,*}, Felix Klempt^a, Hendrik Geisler^a, Rumjhum
5 Mukherjee^{b,c}, Nils Heine^{b,c}, Katharina Doll-Nikutta^{b,c}, Meike Stiesch^{b,c},
6 Szymon P. Szafrński^{b,c}, Meisam Soleimani^a, Philipp Junker^a

7 ^a*Institute of Continuum Mechanics, Leibniz Universität Hannover, An der Universität 1,
8 30823 Garbsen, Germany*

9 ^b*Department of Prosthetic Dentistry and Biomedical Materials Science, Hannover
10 Medical School, Hannover, Germany*

11 ^c*Lower Saxony Centre for Biomedical Engineering, Implant Research and Development
12 (NIFE), Hannover, Germany*

13 ^d*School of Science for Open and Environmental Systems, Graduate School of Science and
14 Technology, Keio University, Yokohama, Japan*

15 **Abstract**

Biofilms are living materials whose composition, and with it their mechanical response, is set by interactions between species that cannot be measured directly. We treat the identification of these interactions as a physics-constrained inverse characterization problem and calibrate a thermodynamically consistent multiphase model of a five-species oral biofilm, derived from the extended Hamilton principle, against in-vitro time courses under four combinations of community state and cultivation. All 15 entries of the symmetric interaction matrix are inferred jointly by transitional Markov chain Monte Carlo, with the stiff forward model compiled in JAX and evaluated for all particles in parallel on a GPU ($\sim 150\times$ faster than on a CPU). The data enter in stages, from species composition to community viability, and every stage is accepted only after explicit checks of seed agreement, mixing and posterior mass at the prior bounds. **[Results: fit, which interactions are identified, the Fn–Pg coefficient, pH check, knockout prediction — after the final runs.]**

16 *Keywords:* Inverse characterization, Physics-constrained Bayesian
17 inference, Scientific machine learning, Extended Hamilton principle, Living

*Corresponding author
Preprint submitted to *Extreme Mechanics Letters* October 10, 2026
Email address: keisuke.nishioka@stud.uni-hannover.de (Keisuke Nishioka)

19 **1. Introduction**

20 Biofilms are soft, living materials: microbial cells embedded in a self-
21 produced matrix whose composition evolves in time and determines their
22 mechanical response [1]. On dental implants, a shift in the species compo-
23 sition of the biofilm from a commensal towards a dysbiotic community is
24 associated with peri-implantitis [2, 3, 4]. Which interactions between the
25 species drive such a shift is difficult to measure: co-culture experiments re-
26 port how the composition changes over time, not the interactions themselves.

27 Continuum mechanics offers a route from composition to interaction.
28 Based on the extended Hamilton principle [5], Klempt et al. [6] derived a
29 multiphase model in which species volume fractions and viabilities evolve
30 under a free energy that contains a symmetric interaction matrix $\mathbf{A}$; ther-
31 modynamic consistency and the symmetry of $\mathbf{A}$ follow from the variational
32 structure. Calibrating $\mathbf{A}$ against data is an inverse characterization prob-
33 lem: the physics fixes the structure of the model, and the data determine its
34 coefficients. Data-driven methods for such problems are most reliable when
35 they are guided by physics and by explicit checks of what the data can de-
36 termine. Once calibrated, the model can be used as an in-silico test bed,
37 for example to ask how the community would develop if one species were
38 removed. Bayesian updating has been used to calibrate constitutive param-
39 eters of biofilm models with surrogates [7], but the interaction matrix poses
40 two specific difficulties. First, with five usable time points and five species,
41 at most 25 observations constrain 15 unknowns, so identifiability cannot be
42 assumed. Second, the coupled volume-fraction and viability dynamics pro-
43 duce posteriors that are multimodal and bounded by the prior box, for which
44 the convergence of standard samplers is hard to verify.

45 Here we infer all 15 entries of $\mathbf{A}$ from the time courses of Heine et al. [8]
46 for five oral taxa—*Streptococcus oralis* (So), *Actinomyces naeslundii* (An),
47 *Veillonella* spp. (Vei), *Fusobacterium nucleatum* (Fn) and *Porphyromonas*
48 *gingivalis* (Pg)—in four conditions: commensal or dysbiotic community, each
49 under static or dynamic (HOBIC flow reactor) cultivation. We combine
50 three elements: (i) transitional Markov chain Monte Carlo (TMCMC) [9, 10]
51 with the forward model compiled in JAX [11] and evaluated for all particles
52 in parallel on a GPU; (ii) a staged introduction of the data, from species

53 composition to community viability; and (iii) explicit acceptance criteria for
 54 every stage, including a rule for widening prior bounds at which posterior
 55 mass accumulates, and an analysis of which interactions the data determine.
 56 **[One sentence on the main result after the final runs.]**

57 2. Model

58 At a material point, species $i = 1, \dots, 5$ occupies the volume fraction
 59 $\phi_i \in (0, 1)$, of which the fraction $\psi_i \in (0, 1)$ is alive; the living volume fraction
 60 is $\bar{\phi}_i = \phi_i \psi_i$. A void fraction ϕ_0 closes the volume, $\sum_{l=0}^5 \phi_l = 1$. Following
 61 Klempt et al. [6], the evolution equations follow from the stationarity of the
 62 rate functional

$$\dot{\Psi} + \Delta + C \rightarrow \text{stat}_{(\phi, \psi, \gamma)}, \quad (1)$$

63 the local form of the extended Hamilton principle [5, 12, 13, 14], with the
 64 free energy, dissipation and constraint

$$\Psi = -\frac{1}{2} c^* \bar{\phi}^\top \mathbf{A} \bar{\phi}, \quad \Delta = \frac{1}{2} \dot{\bar{\phi}}^\top \boldsymbol{\eta} \dot{\bar{\phi}} + \frac{1}{2} \dot{\phi}^\top \boldsymbol{\eta} \dot{\phi}, \quad C = \gamma \left(\sum_{l=0}^5 \phi_l - 1 \right). \quad (2)$$

65 Here c^* is the nutrient concentration, $\boldsymbol{\eta} = \text{diag}(\eta_i)$ with $\eta_i = 1$, and γ
 66 is a Lagrange multiplier; the antibiotic term of the original model vanishes
 67 because the experiments contain no antibiotic. Stationarity gives $2n+2 = 12$
 68 coupled non-linear equations, for $i = 1, \dots, 5$,

$$0 = -c^* \psi_i (\mathbf{A} \bar{\phi})_i + \eta_i (\dot{\phi}_i \psi_i^2 + \bar{\phi}_i \dot{\psi}_i + \dot{\phi}_i) + \gamma, \quad (3)$$

$$0 = -c^* \phi_i (\mathbf{A} \bar{\phi})_i + \eta_i (\dot{\psi}_i \phi_i^2 + \bar{\phi}_i \dot{\phi}_i) + \gamma, \quad (4)$$

69 together with $0 = \gamma + \dot{\phi}_0$ and the volume constraint. They are integrated
 70 by an implicit Euler scheme with Newton iterations (2500 steps over the 21-
 71 day window), with a logarithmic barrier keeping ϕ_i, ψ_i inside $(0, 1)$. Because
 72 the free energy is quadratic in $\bar{\phi}$, $\mathbf{A}$ is symmetric and has 15 independent
 73 entries, which form the parameter vector $\boldsymbol{\theta} \in \mathbb{R}^{15}$. The coefficient a_{ij} is an
 74 effective, undirected interaction of the model; it lumps mechanisms that the
 75 model does not resolve, such as metabolite exchange or co-aggregation. Only
 76 the product $c^* a_{ij}$ enters, so we fix $c^* = 25$ to keep $a_{ij} = \mathcal{O}(1)$. Removing a
 77 species is represented by holding its volume fraction at zero.

3. Bayesian inference

Likelihood. The data are the species fractions at Days 3, 6, 10, 15 and 21 (Day 1 is the initial condition), averaged over eight biological replicates [8]. The model fractions are renormalised over the five species, $\hat{y}_i = \phi_i / \sum_j \phi_j$, and compared through a weighted Gaussian log-likelihood

$$\ln p(\mathbf{y}_{\text{obs}} | \boldsymbol{\theta}) = -\frac{1}{2} \sum_k \sum_{i=1}^5 w_{ik} \frac{(y_{\text{obs},i}(t_k) - \hat{y}_i(t_k; \boldsymbol{\theta}))^2}{\sigma_i^2} + \text{const}, \quad (5)$$

with σ_i the replicate standard deviation of species i . The weights up-weight Pg ($\lambda_{\text{Pg}} = 5$), whose late increase is the main dynamic feature of the dysbiotic data, and the last two time points ($\lambda_{\text{late}} = 3$), and down-weight species below 5% (0.1); without these weights the median of the key coefficient is unchanged but its interval widens (Supplementary Section S1). In the later stages a second Gaussian channel compares the predicted community viability $\sum_i \hat{\phi}_i / \sum_i \phi_i$ with the measured fraction of membrane-intact cells. The measured pH under dynamic cultivation is not used in the inference and serves as a check (Section 4).

Prior and stages. The prior is uniform on a box $[l_k, u_k]$ for each entry, set to contain earlier estimates; it encodes no sign constraint. Inference proceeds in four stages on the same data: (i) composition likelihood on the box; (ii) as (i), on the box narrowed to the stage-(i) posterior mean ± 4 standard deviations; (iii) composition and viability, back on the full box, starting from the stage-(ii) particles; (iv) as (iii), on the box narrowed to the stage-(iii) posterior mean ± 4 standard deviations. A ± 4 SD box removes less than 10^{-4} of a Gaussian marginal, so narrowing shortens the tempering without truncating the posterior. If more than 20% of the posterior mass of an entry lies in the outer 5% of the box on one side, the bound is moved outwards and the stage is repeated; the boxes finally used are listed in Supplementary Table S1.

TMCMC on the GPU. TMCMC [9, 15] passes a particle population through the tempered posteriors $p_m(\boldsymbol{\theta}) \propto p(\mathbf{y}_{\text{obs}} | \boldsymbol{\theta})^{\beta_m} p(\boldsymbol{\theta})$, $0 = \beta_0 < \dots < \beta_M = 1$, choosing each increment so that the coefficient of variation of the importance weights stays at a target value [10]. After resampling, every particle performs K Metropolis–Hastings steps with random-walk proposals scaled by the weighted particle covariance, alternating with differential-evolution proposals [16] that follow strongly correlated directions; proposals outside

the box are rejected without solving the model. Each likelihood evaluation requires one solution of Eqs. (3)–(4), and a run requires 10^5 – 10^6 of them. We implement the implicit solver, including Jacobian assembly and linear solve, in JAX, compile it into a single XLA kernel and evaluate all particles at once with `vmap`. On one GPU (NVIDIA RTX 4090) this is 170 times faster per stage than a 12-core CPU at matched settings (147 times end to end), which makes it affordable to repeat every stage with three seeds.

Acceptance criteria. Every stage is run with three independent seeds and accepted only if (a) the stored log-likelihood of every particle agrees with a recomputation; (b) the tempering reaches $\beta = 1$; (c) each particle is moved at least twice per tempering step and at least five times per free parameter over the run; (d) the maximum log-likelihood differs by at most 1 across seeds; and (e) every posterior median differs across seeds by at most half a posterior standard deviation; an entry whose pooled posterior standard deviation is at least 0.8 times that of a uniform distribution on its box is reported as not identified by the data and excluded from this criterion, and for an entry with several posterior modes the criterion is applied within each mode, with the mode weights reported together with their spread across seeds. Criterion (c) counts accepted moves rather than the acceptance rate: with optimally scaled random-walk proposals [17], $5d$ accepted moves displace each coordinate by about five posterior standard deviations, independently of the dimension d . Stages that fail are repeated with more particles or mutation steps, never interpreted. The evidence estimate of TMCMC is used only to compare models on the same data and box (Section 4).

4. Results and discussion

[After the final runs (budget 1300 + 550 words, figures 2–6): fits in four conditions (Fig. 2); inferred matrices (Fig. 3) and posteriors with identified/unidentified entries (Fig. 4); the Fn–Pg coefficient: stage (i)–(iv), Bayes factor, likelihood weights; identifiability: wide box with and without a Gaussian prior, Phase 1 vs. Phase 2 (Fig. 5); pH check with a fixed regression on species fractions and the Fn–knockout prediction (Fig. 6); limitations: unit-sum constraint, effective parameters; outlook: propagation to mechanics (stiffness of the biofilm from its composition).]

144 5. Conclusions

145 [After the final runs.]

146 Declaration of competing interest

147 The authors declare no competing interests.

148 Data and code availability

149 Code, run configurations, convergence records and posterior samples are
150 archived at Zenodo ([DOI]) and developed at [https://github.com/keisuke58/](https://github.com/keisuke58/Tmcmc202601)
151 Tmcmc202601. The experimental data are those of Heine et al. [8].

152 Funding

153 This work was funded by the Deutsche Forschungsgemeinschaft (DFG,
154 German Research Foundation) under the Collaborative Research Center SFB/TRR-
155 298-SIIRI – Project ID 426335750. Funded by the European Union (ERC,
156 Gen-TSM, project number 101124463). Views and opinions expressed are
157 however those of the author(s) only and do not necessarily reflect those of
158 the European Union or the European Research Council Executive Agency.
159 Neither the European Union nor the granting authority can be held respon-
160 sible for them.

161 Acknowledgements

162 The authors acknowledge the Lower Saxony Centre for Biomedical Engi-
163 neering, Implant Research and Development (NIFE), Hannover. Computa-
164 tional resources of the Multiphysics Materials Computation Research Group
165 (Assoc. Prof. M. Muramatsu), Department of Mechanical Engineering, Keio
166 University, are gratefully acknowledged.

167 CRediT authorship contribution statement

168 [draft, to be confirmed by all co-authors] **Keisuke Nishioka:** Con-
169 ceptualization, Methodology, Software, Formal analysis, Investigation, Visu-
170 alization, Writing – original draft. **Felix Klemp:** Methodology, Software,
171 Writing – review & editing. **Hendrik Geisler:** Methodology, Writing –

review & editing. **Rumjhum Mukherjee:** Investigation, Writing – review
& editing. **Nils Heine:** Investigation, Data curation, Writing – review &
editing. **Katharina Doll-Nikutta:** Resources, Data curation, Writing –
review & editing. **Meike Stiesch:** Resources, Writing – review & edit-
ing. **Szymon P. Szafrński:** Investigation, Writing – review & editing.
Meisam Soleimani: Methodology, Supervision, Writing – review & edit-
ing. **Philipp Junker:** Conceptualization, Supervision, Funding acquisition,
Writing – review & editing.

References

- [1] F. Klempt, M. Soleimani, P. Wriggers, P. Junker, A Hamilton
principle-based model for diffusion-driven biofilm growth, *Biomechanics
and Modeling in Mechanobiology* 23 (2024) 2091–2113. doi:10.1007/
s10237-024-01883-x.
- [2] P. E. Kolenbrander, R. J. Palmer Jr., S. Periasamy, N. S. Jakubovics,
Oral multispecies biofilm development and the key role of cell–cell dis-
tance, *Nature Reviews Microbiology* 8 (2010) 471–480.
- [3] P. D. Marsh, Are dental diseases examples of ecological catastrophes?,
Microbiology 149 (2003) 279–294. doi:10.1099/mic.0.26082-0.
- [4] S. Dieckow, S. P. Szafrński, J. Grischke, T. Qu, K. Doll-Nikutta,
M. Steglich, I. Yang, S. Häussler, M. Stiesch, Structure and compo-
sition of early biofilms formed on dental implants are complex, diverse,
subject-specific and dynamic, *npj Biofilms and Microbiomes* 10 (2024)
155. doi:10.1038/s41522-024-00624-3.
- [5] P. Junker, D. Balzani, An extended Hamilton principle as unify-
ing theory for coupled problems and dissipative microstructure evolu-
tion, *Continuum Mechanics and Thermodynamics* 33 (2021) 1931–1956.
doi:10.1007/s00161-021-01017-z.
- [6] F. Klempt, H. Geisler, M. Soleimani, P. Junker, A contin-
uum multi-species bacterial growth model with a novel interaction
scheme, *Archive of Applied Mechanics* 96 (2026) 164. doi:10.1007/
s00419-026-03160-y.

- [7] L. Fritsch, H. Geisler, J. Grashorn, F. Klempt, M. Soleimani, M. Broggi, P. Junker, M. Beer, Bayesian Updating of constitutive parameters under hybrid uncertainties with a novel surrogate model applied to biofilms, *Computational Mechanics* 78 (2026) 761–782. doi:10.1007/s00466-026-02777-8.
- [8] N. Heine, K. Bittroff, S. P. Szafranski, M. Duitscher, W. Behrens, C. Vollmer, C. Mikolai, N. Kommerein, N. Debener, K. Frings, A. Heisterkamp, T. Scheper, M. L. Torres-Mapa, J. Bahnemann, M. Stiesch, K. Doll-Nikutta, Influence of species composition and cultivation condition on peri-implant biofilm dysbiosis in vitro, *Frontiers in Oral Health* 6 (2025) 1649419. doi:10.3389/froh.2025.1649419.
- [9] J. Ching, Y.-C. Chen, Transitional Markov chain Monte Carlo method for Bayesian model updating, model class selection, and model averaging, *Journal of Engineering Mechanics* 133 (2007) 816–832.
- [10] W. Betz, I. Papaioannou, D. Straub, Transitional Markov chain Monte Carlo: Observations and improvements, *Journal of Engineering Mechanics* 142 (2016) 04016016.
- [11] J. Bradbury, R. Frostig, P. Hawkins, M. J. Johnson, C. Leary, D. Maclaurin, G. Necula, A. Paszke, J. VanderPlas, S. Wanderman-Milne, Q. Zhang, JAX: Composable transformations of Python+NumPy programs, 2018. URL: <https://github.com/google/jax>.
- [12] P. Junker, T. Bode, K. Hackl, On a holistic variational formulation for material modeling including dissipative evolution, *Journal of the Mechanics and Physics of Solids* 200 (2025) 106133. doi:10.1016/j.jmps.2025.106133.
- [13] P. Junker, J. Makowski, K. Hackl, The principle of the minimum of the dissipation potential for non-isothermal processes, *Continuum Mechanics and Thermodynamics* 26 (2014) 259–268. doi:10.1007/s00161-013-0299-4.
- [14] K. Hackl, Generalized standard media and variational principles in classical and finite elastoplasticity, *Journal of the Mechanics and Physics of Solids* 45 (1997) 667–688. doi:10.1016/S0022-5096(96)00110-X.

- 235 [15] P. Del Moral, A. Doucet, A. Jasra, Sequential Monte Carlo samplers,
236 Journal of the Royal Statistical Society: Series B 68 (2006) 411–436.
- 237 [16] C. J. F. ter Braak, A Markov chain Monte Carlo version of the ge-
238 netic algorithm Differential Evolution: easy Bayesian computing for
239 real parameter spaces, Statistics and Computing 16 (2006) 239–249.
240 doi:10.1007/s11222-006-8769-1.
- 241 [17] G. O. Roberts, J. S. Rosenthal, Optimal scaling for various Metropolis–
242 Hastings algorithms, Statistical Science 16 (2001) 351–367. doi:10.
243 1214/ss/1015346320.