

GPU-accelerated Bayesian calibration of a variational
multiphase model of multispecies biofilms **[final title
after the results]**

Keisuke Nishioka^{a,d,*}, Felix Klempt^a, Hendrik Geisler^a, Rumjhum
Mukherjee^{b,c}, Nils Heine^{b,c}, Katharina Doll-Nikutta^{b,c}, Meike Stiesch^{b,c},
Szymon P. Szafrński^{b,c}, Meisam Soleimani^a, Philipp Junker^a

^a*Institute of Continuum Mechanics, Leibniz Universität Hannover, An der Universität 1,
30823 Garbsen, Germany*

^b*Department of Prosthetic Dentistry and Biomedical Materials Science, Hannover
Medical School, Hannover, Germany*

^c*Lower Saxony Centre for Biomedical Engineering, Implant Research and Development
(NIFE), Hannover, Germany*

^d*School of Science for Open and Environmental Systems, Graduate School of Science and
Technology, Keio University, Yokohama, Japan*

Abstract

Biofilms are living materials whose composition, and with it their mechanical response, is set by interactions between species that cannot be measured directly. We treat the identification of these interactions as a physics-constrained inverse characterization problem and calibrate a thermodynamically consistent multiphase model of a five-species oral biofilm, derived from the extended Hamilton principle, against in-vitro time courses under four combinations of community state and cultivation. All 15 entries of the symmetric interaction matrix are inferred jointly by transitional Markov chain Monte Carlo, with the stiff forward model compiled in JAX and evaluated for all particles in parallel on a GPU ($\sim 150\times$ faster than on a CPU). The data enter in stages, from species composition to community viability, and every stage is accepted only after explicit checks of seed agreement, mixing and posterior mass at the prior bounds. **[Results: fit, which interactions are identified, the Fn–Pg coefficient, pH check, knockout prediction — after the final runs.]**

Keywords: Inverse characterization, Physics-constrained Bayesian inference, Scientific machine learning, Extended Hamilton principle, Living

*Corresponding author
Preprint submitted to *Extreme Mechanics Letters* October 9, 2026
Email address: keisuke.nishioka@stud.uni-hannover.de (Keisuke Nishioka)

19 **1. Introduction**

20 Biofilms are soft, living materials: microbial cells embedded in a self-
21 produced matrix whose composition evolves in time and determines their
22 mechanical response [1]. On dental implants, a shift in the species compo-
23 sition of the biofilm from a commensal towards a dysbiotic community is
24 associated with peri-implantitis [2, 3, 4]. Which interactions between the
25 species drive such a shift is difficult to measure: co-culture experiments re-
26 port how the composition changes over time, not the interactions themselves.

27 Continuum mechanics offers a route from composition to interaction.
28 Based on the extended Hamilton principle [5], Klempt et al. [6] derived a
29 multiphase model in which species volume fractions and viabilities evolve
30 under a free energy that contains a symmetric interaction matrix $\mathbf{A}$; ther-
31 modynamic consistency and the symmetry of $\mathbf{A}$ follow from the variational
32 structure. Calibrating $\mathbf{A}$ against data is an inverse characterization prob-
33 lem: the physics fixes the structure of the model, and the data determine its
34 coefficients. Data-driven methods for such problems are most reliable when
35 they are guided by physics and by explicit checks of what the data can de-
36 termine. Once calibrated, the model can be used as an in-silico test bed,
37 for example to ask how the community would develop if one species were
38 removed. Bayesian updating has been used to calibrate constitutive param-
39 eters of biofilm models with surrogates [7], but the interaction matrix poses
40 two specific difficulties. First, with five usable time points and five species,
41 at most 25 observations constrain 15 unknowns, so identifiability cannot be
42 assumed. Second, the coupled volume-fraction and viability dynamics pro-
43 duce posteriors that are multimodal and bounded by the prior box, for which
44 the convergence of standard samplers is hard to verify.

45 Here we infer all 15 entries of $\mathbf{A}$ from the time courses of Heine et al. [8]
46 for five oral taxa—*Streptococcus oralis* (So), *Actinomyces naeslundii* (An),
47 *Veillonella* spp. (Vei), *Fusobacterium nucleatum* (Fn) and *Porphyromonas*
48 *gingivalis* (Pg)—in four conditions: commensal or dysbiotic community, each
49 under static or dynamic (HOBIC flow reactor) cultivation. We combine
50 three elements: (i) transitional Markov chain Monte Carlo (TMCMC) [9, 10]
51 with the forward model compiled in JAX [11] and evaluated for all particles
52 in parallel on a GPU; (ii) a staged introduction of the data, from species

composition to community viability; and (iii) explicit acceptance criteria for every stage, including a rule for widening prior bounds at which posterior mass accumulates, and an analysis of which interactions the data determine. **[One sentence on the main result after the final runs.]**

2. Model

At a material point, species $i = 1, \dots, 5$ occupies the volume fraction $\phi_i \in (0, 1)$, of which the fraction $\psi_i \in (0, 1)$ is alive; the living volume fraction is $\bar{\phi}_i = \phi_i \psi_i$. A void fraction ϕ_0 closes the volume, $\sum_{l=0}^5 \phi_l = 1$. Following Klempt et al. [6], the evolution equations follow from the stationarity of the rate functional

$$\dot{\Psi} + \Delta + C \rightarrow \text{stat}_{(\phi, \psi, \gamma)}, \quad (1)$$

the local form of the extended Hamilton principle [5, 12, 13, 14], with the free energy, dissipation and constraint

$$\Psi = -\frac{1}{2} c^* \bar{\phi}^\top \mathbf{A} \bar{\phi}, \quad \Delta = \frac{1}{2} \dot{\bar{\phi}}^\top \boldsymbol{\eta} \dot{\bar{\phi}} + \frac{1}{2} \dot{\phi}^\top \boldsymbol{\eta} \dot{\phi}, \quad C = \gamma \left(\sum_{l=0}^5 \phi_l - 1 \right). \quad (2)$$

Here c^* is the nutrient concentration, $\boldsymbol{\eta} = \text{diag}(\eta_i)$ with $\eta_i = 1$, and γ is a Lagrange multiplier; the antibiotic term of the original model vanishes because the experiments contain no antibiotic. Stationarity gives $2n+2 = 12$ coupled non-linear equations, for $i = 1, \dots, 5$,

$$0 = -c^* \psi_i (\mathbf{A} \bar{\phi})_i + \eta_i (\dot{\phi}_i \psi_i^2 + \bar{\phi}_i \dot{\psi}_i + \dot{\phi}_i) + \gamma, \quad (3)$$

$$0 = -c^* \phi_i (\mathbf{A} \bar{\phi})_i + \eta_i (\dot{\psi}_i \phi_i^2 + \bar{\phi}_i \dot{\phi}_i) + \gamma, \quad (4)$$

together with $0 = \gamma + \dot{\phi}_0$ and the volume constraint. They are integrated by an implicit Euler scheme with Newton iterations (2500 steps over the 21-day window), with a logarithmic barrier keeping ϕ_i, ψ_i inside $(0, 1)$. Because the free energy is quadratic in $\bar{\phi}$, $\mathbf{A}$ is symmetric and has 15 independent entries, which form the parameter vector $\boldsymbol{\theta} \in \mathbb{R}^{15}$. The coefficient a_{ij} is an effective, undirected interaction of the model; it lumps mechanisms that the model does not resolve, such as metabolite exchange or co-aggregation. Only the product $c^* a_{ij}$ enters, so we fix $c^* = 25$ to keep $a_{ij} = \mathcal{O}(1)$. Removing a species is represented by holding its volume fraction at zero.

3. Bayesian inference

Likelihood. The data are the species fractions at Days 3, 6, 10, 15 and 21 (Day 1 is the initial condition), averaged over eight biological replicates [8]. The model fractions are renormalised over the five species, $\hat{y}_i = \phi_i / \sum_j \phi_j$, and compared through a weighted Gaussian log-likelihood

$$\ln p(\mathbf{y}_{\text{obs}} | \boldsymbol{\theta}) = -\frac{1}{2} \sum_k \sum_{i=1}^5 w_{ik} \frac{(y_{\text{obs},i}(t_k) - \hat{y}_i(t_k; \boldsymbol{\theta}))^2}{\sigma_i^2} + \text{const}, \quad (5)$$

with σ_i the replicate standard deviation of species i . The weights up-weight Pg ($\lambda_{\text{Pg}} = 5$), whose late increase is the main dynamic feature of the dysbiotic data, and the last two time points ($\lambda_{\text{late}} = 3$), and down-weight species below 5% (0.1); without these weights the median of the key coefficient is unchanged but its interval widens (Supplementary Section S1). In the later stages a second Gaussian channel compares the predicted community viability $\sum_i \hat{\phi}_i / \sum_i \phi_i$ with the measured fraction of membrane-intact cells. The measured pH under dynamic cultivation is not used in the inference and serves as a check (Section 4).

Prior and stages. The prior is uniform on a box $[l_k, u_k]$ for each entry, set to contain earlier estimates; it encodes no sign constraint. Inference proceeds in four stages on the same data: (i) composition likelihood on the box; (ii) as (i), on the box narrowed to the stage-(i) posterior mean ± 4 standard deviations; (iii) composition and viability, back on the full box, starting from the stage-(ii) particles; (iv) as (iii), on the box narrowed to the stage-(iii) posterior mean ± 4 standard deviations. A ± 4 SD box removes less than 10^{-4} of a Gaussian marginal, so narrowing shortens the tempering without truncating the posterior. If more than 20% of the posterior mass of an entry lies in the outer 5% of the box on one side, the bound is moved outwards and the stage is repeated; the boxes finally used are listed in Supplementary Table S1.

TMCMC on the GPU. TMCMC [9, 15] passes a particle population through the tempered posteriors $p_m(\boldsymbol{\theta}) \propto p(\mathbf{y}_{\text{obs}} | \boldsymbol{\theta})^{\beta_m} p(\boldsymbol{\theta})$, $0 = \beta_0 < \dots < \beta_M = 1$, choosing each increment so that the coefficient of variation of the importance weights stays at a target value [10]. After resampling, every particle performs K Metropolis–Hastings steps with random-walk proposals scaled by the weighted particle covariance, alternating with differential-evolution proposals [16] that follow strongly correlated directions; proposals outside

the box are rejected without solving the model. Each likelihood evaluation requires one solution of Eqs. (3)–(4), and a run requires 10^5 – 10^6 of them. We implement the implicit solver, including Jacobian assembly and linear solve, in JAX, compile it into a single XLA kernel and evaluate all particles at once with `vmap`. On one GPU (NVIDIA RTX 4090) this is 170 times faster per stage than a 12-core CPU at matched settings (147 times end to end), which makes it affordable to repeat every stage with three seeds.

Acceptance criteria. Every stage is run with three independent seeds and accepted only if (a) the stored log-likelihood of every particle agrees with a recomputation; (b) the tempering reaches $\beta = 1$; (c) each particle is moved at least twice per tempering step and at least five times per free parameter over the run; (d) the maximum log-likelihood differs by at most 1 across seeds; and (e) every posterior median differs across seeds by at most half a posterior standard deviation. Criterion (c) counts accepted moves rather than the acceptance rate: with optimally scaled random-walk proposals [17], $5d$ accepted moves displace each coordinate by about five posterior standard deviations, independently of the dimension d . Stages that fail are repeated with more particles or mutation steps, never interpreted. The evidence estimate of TMCMC is used only to compare models on the same data and box (Section 4).

4. Results and discussion

[After the final runs (budget 1300 + 550 words, figures 2–6): fits in four conditions (Fig. 2); inferred matrices (Fig. 3) and posteriors with identified/unidentified entries (Fig. 4); the Fn–Pg coefficient: stage (i)–(iv), Bayes factor, likelihood weights; identifiability: wide box with and without a Gaussian prior, Phase 1 vs. Phase 2 (Fig. 5); pH check with a fixed regression on species fractions and the Fn-knockout prediction (Fig. 6); limitations: unit-sum constraint, effective parameters; outlook: propagation to mechanics (stiffness of the biofilm from its composition).]

5. Conclusions

[After the final runs.]

142 Declaration of competing interest

143 The authors declare no competing interests.

144 Data and code availability

145 Code, run configurations, convergence records and posterior samples are
146 archived at Zenodo ([\[DOI\]](#)) and developed at [https://github.com/keisuke58/](https://github.com/keisuke58/Tmcmc202601)
147 `Tmcmc202601`. The experimental data are those of Heine et al. [8].

148 Funding

149 This work was funded by the Deutsche Forschungsgemeinschaft (DFG,
150 German Research Foundation) under the Collaborative Research Center SFB/TRR-
151 298-SIIRI – Project ID 426335750. Funded by the European Union (ERC,
152 Gen-TSM, project number 101124463). Views and opinions expressed are
153 however those of the author(s) only and do not necessarily reflect those of
154 the European Union or the European Research Council Executive Agency.
155 Neither the European Union nor the granting authority can be held respon-
156 sible for them.

157 Acknowledgements

158 The authors acknowledge the Lower Saxony Centre for Biomedical Engi-
159 neering, Implant Research and Development (NIFE), Hannover. Computa-
160 tional resources of the Multiphysics Materials Computation Research Group
161 (Assoc. Prof. M. Muramatsu), Department of Mechanical Engineering, Keio
162 University, are gratefully acknowledged.

163 CRediT authorship contribution statement

164 **[draft, to be confirmed by all co-authors]** **Keisuke Nishioka:** Con-
165 ceptualization, Methodology, Software, Formal analysis, Investigation, Visu-
166 alization, Writing – original draft. **Felix Klempt:** Methodology, Software,
167 Writing – review & editing. **Hendrik Geisler:** Methodology, Writing –
168 review & editing. **Rumjhum Mukherjee:** Investigation, Writing – review
169 & editing. **Nils Heine:** Investigation, Data curation, Writing – review &
170 editing. **Katharina Doll-Nikutta:** Resources, Data curation, Writing –

171 review & editing. **Meike Stiesch:** Resources, Writing – review & edit-
 172 ing. **Szymon P. Szafranski:** Investigation, Writing – review & editing.
 173 **Meisam Soleimani:** Methodology, Supervision, Writing – review & edit-
 174 ing. **Philipp Junker:** Conceptualization, Supervision, Funding acquisition,
 175 Writing – review & editing.

176 References

- 177 [1] F. Klempt, M. Soleimani, P. Wriggers, P. Junker, A Hamilton
 178 principle-based model for diffusion-driven biofilm growth, *Biomechanics*
 179 and *Modeling in Mechanobiology* 23 (2024) 2091–2113. doi:10.1007/
 180 s10237-024-01883-x.
- 181 [2] P. E. Kolenbrander, R. J. Palmer Jr., S. Periasamy, N. S. Jakubovics,
 182 Oral multispecies biofilm development and the key role of cell–cell dis-
 183 tance, *Nature Reviews Microbiology* 8 (2010) 471–480.
- 184 [3] P. D. Marsh, Are dental diseases examples of ecological catastrophes?,
 185 *Microbiology* 149 (2003) 279–294. doi:10.1099/mic.0.26082-0.
- 186 [4] S. Dieckow, S. P. Szafranski, J. Grischke, T. Qu, K. Doll-Nikutta,
 187 M. Steglich, I. Yang, S. Häussler, M. Stiesch, Structure and compo-
 188 sition of early biofilms formed on dental implants are complex, diverse,
 189 subject-specific and dynamic, *npj Biofilms and Microbiomes* 10 (2024)
 190 155. doi:10.1038/s41522-024-00624-3.
- 191 [5] P. Junker, D. Balzani, An extended Hamilton principle as unify-
 192 ing theory for coupled problems and dissipative microstructure evolu-
 193 tion, *Continuum Mechanics and Thermodynamics* 33 (2021) 1931–1956.
 194 doi:10.1007/s00161-021-01017-z.
- 195 [6] F. Klempt, H. Geisler, M. Soleimani, P. Junker, A contin-
 196 uum multi-species bacterial growth model with a novel interaction
 197 scheme, *Archive of Applied Mechanics* 96 (2026) 164. doi:10.1007/
 198 s00419-026-03160-y.
- 199 [7] L. Fritsch, H. Geisler, J. Grashorn, F. Klempt, M. Soleimani, M. Broggi,
 200 P. Junker, M. Beer, Bayesian Updating of constitutive parameters
 201 under hybrid uncertainties with a novel surrogate model applied to
 202 biofilms, *Computational Mechanics* 78 (2026) 761–782. doi:10.1007/
 203 s00466-026-02777-8.

- [8] N. Heine, K. Bittroff, S. P. Szafranski, M. Duitscher, W. Behrens, C. Vollmer, C. Mikolai, N. Kommerein, N. Debener, K. Frings, A. Heisterkamp, T. Scheper, M. L. Torres-Mapa, J. Bahnemann, M. Stiesch, K. Doll-Nikutta, Influence of species composition and cultivation condition on peri-implant biofilm dysbiosis in vitro, *Frontiers in Oral Health* 6 (2025) 1649419. doi:10.3389/froh.2025.1649419.
- [9] J. Ching, Y.-C. Chen, Transitional Markov chain Monte Carlo method for Bayesian model updating, model class selection, and model averaging, *Journal of Engineering Mechanics* 133 (2007) 816–832.
- [10] W. Betz, I. Papaioannou, D. Straub, Transitional Markov chain Monte Carlo: Observations and improvements, *Journal of Engineering Mechanics* 142 (2016) 04016016.
- [11] J. Bradbury, R. Frostig, P. Hawkins, M. J. Johnson, C. Leary, D. Maclaurin, G. Necula, A. Paszke, J. VanderPlas, S. Wanderman-Milne, Q. Zhang, JAX: Composable transformations of Python+NumPy programs, 2018. URL: <https://github.com/google/jax>.
- [12] P. Junker, T. Bode, K. Hackl, On a holistic variational formulation for material modeling including dissipative evolution, *Journal of the Mechanics and Physics of Solids* 200 (2025) 106133. doi:10.1016/j.jmps.2025.106133.
- [13] P. Junker, J. Makowski, K. Hackl, The principle of the minimum of the dissipation potential for non-isothermal processes, *Continuum Mechanics and Thermodynamics* 26 (2014) 259–268. doi:10.1007/s00161-013-0299-4.
- [14] K. Hackl, Generalized standard media and variational principles in classical and finite elastoplasticity, *Journal of the Mechanics and Physics of Solids* 45 (1997) 667–688. doi:10.1016/S0022-5096(96)00110-X.
- [15] P. Del Moral, A. Doucet, A. Jasra, Sequential Monte Carlo samplers, *Journal of the Royal Statistical Society: Series B* 68 (2006) 411–436.
- [16] C. J. F. ter Braak, A Markov chain Monte Carlo version of the genetic algorithm Differential Evolution: easy Bayesian computing for real parameter spaces, *Statistics and Computing* 16 (2006) 239–249. doi:10.1007/s11222-006-8769-1.

- 237 [17] G. O. Roberts, J. S. Rosenthal, Optimal scaling for various Metropolis–
238 Hastings algorithms, *Statistical Science* 16 (2001) 351–367. doi:10.
239 1214/ss/1015346320.