GPU-accelerated Bayesian calibration of a variational multiphase model of multispecies biofilms [**final title after the results**]

Keisuke Nishioka[a,d,*], Felix Klempt[a], Hendrik Geisler[a], Rumjhum Mukherjee[b,c], Nils Heine[b,c], Katharina Doll-Nikutta[b,c], Meike Stiesch[b,c], Szymon P. Szafrański[b,c], Meisam Soleimani[a], Philipp Junker[a]

[a] *Institute of Continuum Mechanics, Leibniz Universität Hannover, An der Universität 1, 30823 Garbsen, Germany*
[b] *Department of Prosthetic Dentistry and Biomedical Materials Science, Hannover Medical School, Hannover, Germany*
[c] *Lower Saxony Centre for Biomedical Engineering, Implant Research and Development (NIFE), Hannover, Germany*
[d] *School of Science for Open and Environmental Systems, Graduate School of Science and Technology, Keio University, Yokohama, Japan*

Abstract

Biofilms are living materials whose composition, and with it their mechanical response, is set by interactions between species that cannot be measured directly. We treat the identification of these interactions as a physics-constrained inverse characterization problem and calibrate a thermodynamically consistent multiphase model of a five-species oral biofilm, derived from the extended Hamilton principle, against in-vitro time courses under four combinations of community state and cultivation. All 15 entries of the symmetric interaction matrix are inferred jointly by transitional Markov chain Monte Carlo, with the stiff forward model compiled in JAX and evaluated for all particles in parallel on a GPU (∼150× faster than on a CPU). The data enter in stages, from species composition to community viability, and every stage is accepted only after explicit checks of seed agreement, mixing and posterior mass at the prior bounds. The calibrated model reproduces the species fractions with root-mean-square errors of 0.03–0.08 in three conditions[**prov.**]; in the commensal static condition (0.18) the composition and the measured viability cannot both be matched. Only part of the interaction matrix is determined by the data—four and six entries in the two commensal conditions, all among the early colonisers and with the same signs where

Preprint submitted to Extreme Mechanics Letters *October 12, 2026*

*Corresponding author
Email address: keisuke.nishioka@stud.uni-hannover.de (Keisuke Nishioka)

they overlap, and eight and nine in the dysbiotic ones[prov.]—and we state explicitly which. The Fn–Pg coefficient, positive when only the composition is used, is no longer sign-determined once the viability is added[prov.]. The measured pH, not used in the inference, is reproduced in the dysbiotic flow condition through a fixed regression on the predicted composition, and removing Fn from the calibrated model abolishes the late increase of Pg, a prediction open to direct experimental test[prov.].

Keywords: Inverse characterization, Physics-constrained Bayesian inference, Scientific machine learning, Extended Hamilton principle, Living materials, GPU

1. Introduction

Biofilms are soft, living materials: microbial cells embedded in a self-produced matrix whose composition evolves in time and determines their mechanical response [1]. On dental implants, a shift in the species composition of the biofilm from a commensal towards a dysbiotic community is associated with peri-implantitis [2, 3, 4]. Which interactions between the species drive such a shift is difficult to measure: co-culture experiments report how the composition changes over time, not the interactions themselves.

Continuum mechanics offers a route from composition to interaction. Based on the extended Hamilton principle [5], Klempt et al. [6] derived a multiphase model in which species volume fractions and viabilities evolve under a free energy that contains a symmetric interaction matrix $\mathbf{A}$; thermodynamic consistency and the symmetry of $\mathbf{A}$ follow from the variational structure. Calibrating $\mathbf{A}$ against data is an inverse characterization problem: the physics fixes the structure of the model, and the data determine its coefficients. Data-driven methods for such problems are most reliable when they are guided by physics and by explicit checks of what the data can determine. Once calibrated, the model can be used as an in-silico test bed, for example to ask how the community would develop if one species were removed. Bayesian updating has been used to calibrate constitutive parameters of biofilm models with surrogates [7], but the interaction matrix poses two specific difficulties. First, with five usable time points and five species, at most 25 observations constrain 15 unknowns, so identifiability cannot be assumed. Second, the coupled volume-fraction and viability dynamics produce posteriors that are multimodal and bounded by the prior box, for which

the convergence of standard samplers is hard to verify.

Here we infer all 15 entries of $\mathbf{A}$ from the time courses of Heine et al. [8] for five oral taxa—*Streptococcus oralis* (So), *Actinomyces naeslundii* (An), *Veillonella* spp. (Vei), *Fusobacterium nucleatum* (Fn) and *Porphyromonas gingivalis* (Pg)—in four conditions: commensal or dysbiotic community, each under static or dynamic (HOBIC flow reactor) cultivation. We combine three elements: (i) transitional Markov chain Monte Carlo (TMCMC) [9, 10] with the forward model compiled in JAX [11] and evaluated for all particles in parallel on a GPU; (ii) a staged introduction of the data, from species composition to community viability; and (iii) explicit acceptance criteria for every stage, including a rule for widening prior bounds at which posterior mass accumulates, and an analysis of which interactions the data determine. The data determine only a subset of $\mathbf{A}$, which we identify entry by entry; where the commensal conditions overlap, the identified entries agree in sign across cultivation modes, whereas the Vei–Pg coefficient of the dysbiotic community changes sign between static and dynamic cultivation**[prov.]**.

2. Model

At a material point, species $i = 1, \dots, 5$ occupies the volume fraction $\phi_i \in (0, 1)$, of which the fraction $\psi_i \in (0, 1)$ is alive; the living volume fraction is $\bar{\phi}_i = \phi_i \psi_i$. A void fraction ϕ_0 closes the volume, $\sum_{l=0}^{5} \phi_l = 1$. Following Klempt et al. [6], the evolution equations follow from the stationarity of the rate functional

$$\dot{\Psi} + \Delta + C \to \mathrm{stat}_{(\dot{\boldsymbol{\phi}}, \dot{\boldsymbol{\psi}}, \gamma)}, \tag{1}$$

the local form of the extended Hamilton principle [5, 12, 13, 14], with the free energy, dissipation and constraint

$$\Psi = -\tfrac{1}{2}\, c^*\, \bar{\boldsymbol{\phi}}^{\mathsf{T}} \mathbf{A}\, \bar{\boldsymbol{\phi}}, \qquad \Delta = \tfrac{1}{2}\, \dot{\bar{\boldsymbol{\phi}}}^{\mathsf{T}} \boldsymbol{\eta}\, \dot{\bar{\boldsymbol{\phi}}} + \tfrac{1}{2}\, \dot{\boldsymbol{\phi}}^{\mathsf{T}} \boldsymbol{\eta}\, \dot{\boldsymbol{\phi}}, \qquad C = \gamma \Big(\sum_{l=0}^{5} \phi_l - 1 \Big). \tag{2}$$

Here c^* is the nutrient concentration, $\boldsymbol{\eta} = \mathrm{diag}(\eta_i)$ with $\eta_i = 1$, and γ is a Lagrange multiplier; the antibiotic term of the original model vanishes because the experiments contain no antibiotic. Stationarity gives $2n+2 = 12$ coupled non-linear equations, for $i = 1, \dots, 5$,

$$0 = -c^* \psi_i (\mathbf{A} \bar{\boldsymbol{\phi}})_i + \eta_i \big(\dot{\phi}_i \psi_i^2 + \bar{\phi}_i \dot{\psi}_i + \dot{\phi}_i \big) + \gamma, \tag{3}$$

$$0 = -c^* \phi_i (\mathbf{A} \bar{\boldsymbol{\phi}})_i + \eta_i \big(\dot{\psi}_i \phi_i^2 + \bar{\phi}_i \dot{\phi}_i \big) + \gamma, \tag{4}$$

together with $0 = \gamma + \dot{\phi}_0$ and the volume constraint. They are integrated by an implicit Euler scheme with Newton iterations (2500 steps over the 21-day window), with a logarithmic barrier keeping ϕ_i, ψ_i inside $(0, 1)$. Because the free energy is quadratic in $\bar{\boldsymbol{\phi}}$, $\mathbf{A}$ is symmetric and has 15 independent entries, which form the parameter vector $\boldsymbol{\theta} \in \mathbb{R}^{15}$. The coefficient a_{ij} is an effective, undirected interaction of the model; it lumps mechanisms that the model does not resolve, such as metabolite exchange or co-aggregation. Only the product $c^* a_{ij}$ enters, so we fix $c^* = 25$ to keep $a_{ij} = \mathcal{O}(1)$. Removing a species is represented by holding its volume fraction at zero.

3. Bayesian inference

Likelihood.. The data are the species fractions at Days 3, 6, 10, 15 and 21 (Day 1 is the initial condition), averaged over eight biological replicates [8]. The model fractions are renormalised over the five species, $\hat{y}_i = \phi_i / \sum_j \phi_j$, and compared through a weighted Gaussian log-likelihood

$$\ln p(\mathbf{y}_{\text{obs}} \mid \boldsymbol{\theta}) = -\frac{1}{2} \sum_{k} \sum_{i=1}^{5} w_{ik} \, \frac{\big(y_{\text{obs},i}(t_k) - \hat{y}_i(t_k; \boldsymbol{\theta})\big)^2}{\sigma_i^2} + \text{const}, \qquad (5)$$

with σ_i the replicate standard deviation of species i. The weights up-weight Pg ($\lambda_{\text{Pg}} = 5$), whose late increase is the main dynamic feature of the dysbiotic data, and the last two time points ($\lambda_{\text{late}} = 3$), and down-weight species below 5% (0.1); without these weights the median of the key coefficient is unchanged but its interval widens (Supplementary Section S[]). In the later stages a second Gaussian channel compares the predicted community viability $\sum_i \bar{\phi}_i / \sum_i \phi_i$ with the measured fraction of membrane-intact cells. The measured pH under dynamic cultivation is not used in the inference and serves as a check (Section 4).

Prior and stages.. The prior is uniform on a box $[l_k, u_k]$ for each entry, set to contain earlier estimates; it encodes no sign constraint. Inference proceeds in four stages on the same data: (i) composition likelihood on the box; (ii) as (i), on the box narrowed to the stage-(i) posterior mean ± 4 standard deviations; (iii) composition and viability, back on the full box, starting from the stage-(ii) particles; (iv) as (iii), on the box narrowed to the stage-(iii) posterior mean ± 4 standard deviations. A ± 4 SD box removes less than 10^{-4} of a Gaussian marginal, so narrowing shortens the tempering without truncating

the posterior. If more than 20% of the posterior mass of an entry lies in the outer 5% of the box on one side, the bound is moved outwards and the stage is repeated; the boxes finally used are listed in Supplementary Table S1.

TMCMC on the GPU.. TMCMC [9, 15] passes a particle population through the tempered posteriors $p_m(\boldsymbol{\theta}) \propto p(\mathbf{y}_{\mathrm{obs}} \mid \boldsymbol{\theta})^{\beta_m} p(\boldsymbol{\theta})$, $0 = \beta_0 < \cdots < \beta_M = 1$, choosing each increment so that the coefficient of variation of the importance weights stays at a target value [10]. After resampling, every particle performs K Metropolis–Hastings steps with random-walk proposals scaled by the weighted particle covariance, alternating with differential-evolution proposals [16] that follow strongly correlated directions; proposals outside the box are rejected without solving the model. Each likelihood evaluation requires one solution of Eqs. (3)–(4), and a run requires 10^5–10^6 of them. We implement the implicit solver, including Jacobian assembly and linear solve, in JAX, compile it into a single XLA kernel and evaluate all particles at once with `vmap`. On one GPU (NVIDIA RTX 4090) this is 170 times faster per stage than a 12-core CPU at matched settings (147 times end to end), which makes it affordable to repeat every stage with three seeds.

Acceptance criteria.. Every stage is run with three independent seeds and accepted only if (a) the stored log-likelihood of every particle agrees with a recomputation; (b) the tempering reaches $\beta = 1$; (c) each particle is moved at least twice per tempering step and at least five times per free parameter over the run; (d) the maximum log-likelihood differs by at most 1 across seeds; and (e) every posterior median differs across seeds by at most half a posterior standard deviation; an entry whose pooled posterior standard deviation is at least 0.8 times that of a uniform distribution on its box is reported as not identified by the data and excluded from this criterion, and for an entry with several posterior modes the criterion is applied within each mode, with the mode weights reported together with their spread across seeds. Criterion (c) counts accepted moves rather than the acceptance rate: with optimally scaled random-walk proposals [17], $5d$ accepted moves displace each coordinate by about five posterior standard deviations, independently of the dimension d. Stages that fail are repeated with more particles or mutation steps, never interpreted. The evidence estimate of TMCMC is used only to compare models on the same data and box (Section 4).

4. Results and discussion

4.1. Fit and convergence

All stages of all conditions were accepted under criteria (a)–(e) with three seeds each **[DS, DH stage (iv) pending; their numbers below are from stage (ii)]**. At stage (iv) the composition is reproduced with a root-mean-square error of 0.076–0.078 in CH, 0.172–0.179 in CS, 0.021–0.028 in DS and 0.081–0.090 in DH**[prov.]** (Fig. 1). Accepting a stage often required more than one attempt: posterior mass accumulated at a bound of the initial box in every condition and was released by widening that bound (Supplementary Table S1), and the dysbiotic conditions needed up to **[18000 (DS) and 20000 (DH)]** particles before the three seeds agreed (Supplementary Table S2). Without the acceptance criteria, several of these runs would have been read as converged.

4.2. What the data determine

Supplementary Table S4 classifies every entry of $\mathbf{A}$ at stage (iv). In the commensal conditions only a few entries are identified: four in CS (a_{11}, a_{22}, a_{33}, a_{13}) and six in CH (a_{11}, a_{12}, a_{22}, a_{33}, a_{13}, a_{23}), all among the early colonisers So, An and Vei. Where both conditions identify an entry, the sign agrees: the self-interaction of So is negative (a_{11}, 90% intervals $[-4.9, -1.2]$ in CS and $[-2.5, -0.5]$ in CH) and the So–Vei coefficient is positive (a_{13}, $[+2.8, +9.2]$ and $[+0.01, +1.4]$), in line with the lactate cross-feeding between streptococci and *Veillonella* [2], although a_{13} is an effective coefficient and not a measurement of that mechanism. Two further entries of CH, a_{14} and a_{15}, are determined in sign only (90% intervals $[-4.8, -0.2]$ and $[-4.8, -0.3]$), and every other entry involving Fn or Pg spreads over its whole box. This is expected: Fn and Pg stay below 1% of the commensal communities from day 3 on, so the composition data carry little information on their interactions. In the dysbiotic conditions, where Fn and Pg reach 10–25% of the community, more entries are identified: eight in DS and nine in DH, now including interactions of the late colonisers**[prov.]**. The Vei–Pg coefficient a_{35} is identified in both but with opposite signs, positive under static cultivation (DS, $[+0.3, +2.4]$) and negative in the flow reactor (DH, $[-13.4, -1.7]$), and the An–Fn coefficient a_{24} is positive in DH ($[+2.7, +7.1]$). A sign change between cultivation modes is a statement about the effective model under different nutrient and flow conditions, not about the pairwise mechanism itself. The classification is a result in its own right: it states which conclusions

about the interaction network the data can support, and it is obtained from the same runs at no additional cost.

4.3. The Fn–Pg coefficient

The coefficient a_{45} links the bridging organism Fn to the pathogen Pg. Comparing the full model with a model in which $a_{45} = 0$, both inferred at stage (i), gives $\ln B = +0.10$ (CS), $+0.05$ (CH), -2.67 (DS) and $+1.94$ (DH) (Supplementary Table S3). Because the uniform prior penalises the full model by the log-width of the a_{45} box, $\ln B$ is read only within a condition: in DS, where the box had to be widened to $[-15, 20]$, the penalty alone is $\ln 35 \approx 3.6$. In DH, the data favour a non-zero coefficient. With the composition alone (stage (i)) a_{45} is positive in all seeds (90% interval about $[+0.5, +5.5]$); without the likelihood weights its median is unchanged but the interval reaches zero ($[-0.06, +5.6]$; Supplementary Section S5). Once the viability enters, the median of a_{45} stays positive in both dysbiotic conditions ($+2.8$ in DS, $+1.2$ in DH) but the 90% intervals include zero ($[-1.3, +15.1]$ and $[-1.9, +7.5]$)**[prov.]**. The composition data alone thus support a positive Fn–Pg coefficient in DH, and the viability data neither confirm nor contradict it.

4.4. Checks against data not used in the inference

The pH measured under dynamic cultivation was not used in the inference. A linear regression of the measured pH on the measured species fractions (intercept and four fractions, Vei as reference; $R^2 = 0.93$ over 90 time points in CH and DH) was fitted once on the data alone and then applied, unchanged, to the posterior predictive composition. In DH the predicted pH follows the measured rise from about 6.3 to 7.0 ($R^2 = 0.63$, RMSE 0.12)**[prov.]**; in CH it is offset upwards by about 0.13, because the model keeps Fn and Pg near 4–5% where the data show below 1%, a consequence of down-weighting rare species in the likelihood, and these two taxa carry the largest regression coefficients. The calibrated model can also be used in silico: removing Fn from the dysbiotic flow-reactor community abolishes the late increase of Pg. The Day-21/Day-15 ratio of Pg falls from 2.9 (90% interval $[2.3, 3.5]$; measured 2.74) to 1.1 ($[0.3, 1.8]$), and only 2–7% of the posterior samples keep a ratio of 1.5 or more**[prov.]**. Setting only $a_{45} = 0$ removes the increase in most but not all samples (16–52% keep it, depending on the seed), so the prediction rests mainly, though not only, on the Fn–Pg coefficient; it can be tested directly by a co-culture without Fn. Under static

cultivation the model shows no comparable late increase (ratio 1.1), so the knockout is not informative there.

4.5. Limitations

The likelihood treats the five species fractions as independent although they sum to one, so the posterior intervals may be somewhat too narrow. The fit is poorest in CS, almost entirely in So and Vei (per-species errors 0.31 and 0.24, all others below 0.05), where the So fraction rises from 0.18 to 0.70 between day 1 and day 3. The model can follow this rise: fitted to the composition alone it reaches an error of 0.075, but only with a live fraction near 0.97, against the measured 0.65–0.77. The composition and the live/dead staining cannot both be matched in CS, and the posterior settles on a compromise; we report it rather than weight the channels differently for one condition. The compromise is sharpest at the MAP: the bulk of the CS posterior lies closer to the composition data (pointwise posterior median, error 0.125), whereas the MAP, identical in all three seeds, sits where the live fraction is matched best (error 0.179); in the other conditions the two differ by less than 0.02 (Supplementary Fig. [S]). The interaction coefficients are effective quantities of this model and data set; they lump mechanisms that the model does not resolve, such as metabolite exchange and pH, and are not direct measurements of pairwise mechanisms. The calibration covers a 0D material point; the spatial structure of the biofilm and its mechanics are not part of the inference. The calibrated dynamics are, however, the input that continuum growth models of biofilms require [1]: the total volume fraction of the point model can drive the growth variable of a finite-element model, so that differences in growth-induced stress between conditions follow from the calibrated interactions rather than from assumed species-specific growth weights. Because the likelihood uses species fractions only, the absolute amount of biomass is a prediction of the model rather than a calibrated quantity, and its posterior spread has to be propagated into such a coupling.

5. Conclusions

The interaction matrix of a variational multiphase biofilm model was calibrated against time courses of a five-species oral community in four conditions, with a GPU-compiled forward model, a staged likelihood and explicit acceptance criteria for every stage. Three points carry over to other inverse

characterizations of living materials. First, convergence has to be checked, not assumed: posterior mass at the prior bounds and disagreement between seeds were found in every condition and removed only by widening bounds and adding particles. Second, identifiability is a result: with 25 observations for 15 unknowns the data determine only some interactions, and stating which ones bounds every biological conclusion drawn from the model. Third, a misfit can be informative: in the commensal static condition the composition and the viability data cannot be matched at once, which points to a process the model lacks rather than to a failed calibration. In the dysbiotic conditions more interactions are identified, and the Fn–Pg coefficient that the composition data alone suggest to be positive is not sign-determined by the full data set[**prov.**].

Declaration of competing interest

The authors declare no competing interests.

Data and code availability

Code, run configurations, convergence records and posterior samples are archived at Zenodo ([**DOI**]) and developed at `https://github.com/keisuke58/Tmcmc202601`. The experimental data are those of Heine et al. [8].

Funding

This work was funded by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) under the Collaborative Research Center SFB/TRR-298-SIIRI – Project ID 426335750. Funded by the European Union (ERC, Gen-TSM, project number 101124463). Views and opinions expressed are however those of the author(s) only and do not necessarily reflect those of the European Union or the European Research Council Executive Agency. Neither the European Union nor the granting authority can be held responsible for them.

Acknowledgements

The authors acknowledge the Lower Saxony Centre for Biomedical Engineering, Implant Research and Development (NIFE), Hannover. Computational resources of the Multiphysics Materials Computation Research Group

(Assoc. Prof. M. Muramatsu), Department of Mechanical Engineering, Keio University, are gratefully acknowledged.

CRediT authorship contribution statement

[draft, to be confirmed by all co-authors] **Keisuke Nishioka:** Conceptualization, Methodology, Software, Formal analysis, Investigation, Visualization, Writing – original draft. **Felix Klempt:** Methodology, Software, Writing – review & editing. **Hendrik Geisler:** Methodology, Writing – review & editing. **Rumjhum Mukherjee:** Investigation, Writing – review & editing. **Nils Heine:** Investigation, Data curation, Writing – review & editing. **Katharina Doll-Nikutta:** Resources, Data curation, Writing – review & editing. **Meike Stiesch:** Resources, Writing – review & editing. **Szymon P. Szafrański:** Investigation, Writing – review & editing. **Meisam Soleimani:** Methodology, Supervision, Writing – review & editing. **Philipp Junker:** Conceptualization, Supervision, Funding acquisition, Writing – review & editing.

References

[1] F. Klempt, M. Soleimani, P. Wriggers, P. Junker, A Hamilton principle-based model for diffusion-driven biofilm growth, Biomechanics and Modeling in Mechanobiology 23 (2024) 2091–2113. doi:10.1007/s10237-024-01883-x.

[2] P. E. Kolenbrander, R. J. Palmer Jr., S. Periasamy, N. S. Jakubovics, Oral multispecies biofilm development and the key role of cell–cell distance, Nature Reviews Microbiology 8 (2010) 471–480.

[3] P. D. Marsh, Are dental diseases examples of ecological catastrophes?, Microbiology 149 (2003) 279–294. doi:10.1099/mic.0.26082-0.

[4] S. Dieckow, S. P. Szafrański, J. Grischke, T. Qu, K. Doll-Nikutta, M. Steglich, I. Yang, S. Häussler, M. Stiesch, Structure and composition of early biofilms formed on dental implants are complex, diverse, subject-specific and dynamic, npj Biofilms and Microbiomes 10 (2024) 155. doi:10.1038/s41522-024-00624-3.

[5] P. Junker, D. Balzani, An extended Hamilton principle as unifying theory for coupled problems and dissipative microstructure evolution, Continuum Mechanics and Thermodynamics 33 (2021) 1931–1956. doi:10.1007/s00161-021-01017-z.

[6] F. Klempt, H. Geisler, M. Soleimani, P. Junker, A continuum multi-species bacterial growth model with a novel interaction scheme, Archive of Applied Mechanics 96 (2026) 164. doi:10.1007/s00419-026-03160-y.

[7] L. Fritsch, H. Geisler, J. Grashorn, F. Klempt, M. Soleimani, M. Broggi, P. Junker, M. Beer, Bayesian Updating of constitutive parameters under hybrid uncertainties with a novel surrogate model applied to biofilms, Computational Mechanics 78 (2026) 761–782. doi:10.1007/s00466-026-02777-8.

[8] N. Heine, K. Bittroff, S. P. Szafrański, M. Duitscher, W. Behrens, C. Vollmer, C. Mikolai, N. Kommerein, N. Debener, K. Frings, A. Heisterkamp, T. Scheper, M. L. Torres-Mapa, J. Bahnemann, M. Stiesch, K. Doll-Nikutta, Influence of species composition and cultivation condition on peri-implant biofilm dysbiosis in vitro, Frontiers in Oral Health 6 (2025) 1649419. doi:10.3389/froh.2025.1649419.

[9] J. Ching, Y.-C. Chen, Transitional Markov chain Monte Carlo method for Bayesian model updating, model class selection, and model averaging, Journal of Engineering Mechanics 133 (2007) 816–832.

[10] W. Betz, I. Papaioannou, D. Straub, Transitional Markov chain Monte Carlo: Observations and improvements, Journal of Engineering Mechanics 142 (2016) 04016016.

[11] J. Bradbury, R. Frostig, P. Hawkins, M. J. Johnson, C. Leary, D. Maclaurin, G. Necula, A. Paszke, J. VanderPlas, S. Wanderman-Milne, Q. Zhang, JAX: Composable transformations of Python+NumPy programs, 2018. URL: https://github.com/google/jax.

[12] P. Junker, T. Bode, K. Hackl, On a holistic variational formulation for material modeling including dissipative evolution, Journal of the Mechanics and Physics of Solids 200 (2025) 106133. doi:10.1016/j.jmps.2025.106133.

[13] P. Junker, J. Makowski, K. Hackl, The principle of the minimum of the dissipation potential for non-isothermal processes, Continuum Mechanics and Thermodynamics 26 (2014) 259–268. doi:10.1007/s00161-013-0299-4.

[14] K. Hackl, Generalized standard media and variational principles in classical and finite elastoplasticity, Journal of the Mechanics and Physics of Solids 45 (1997) 667–688. doi:10.1016/S0022-5096(96)00110-X.

[15] P. Del Moral, A. Doucet, A. Jasra, Sequential Monte Carlo samplers, Journal of the Royal Statistical Society: Series B 68 (2006) 411–436.

[16] C. J. F. ter Braak, A Markov chain Monte Carlo version of the genetic algorithm Differential Evolution: easy Bayesian computing for real parameter spaces, Statistics and Computing 16 (2006) 239–249. doi:10.1007/s11222-006-8769-1.

[17] G. O. Roberts, J. S. Rosenthal, Optimal scaling for various Metropolis–Hastings algorithms, Statistical Science 16 (2001) 351–367. doi:10.1214/ss/1015346320.

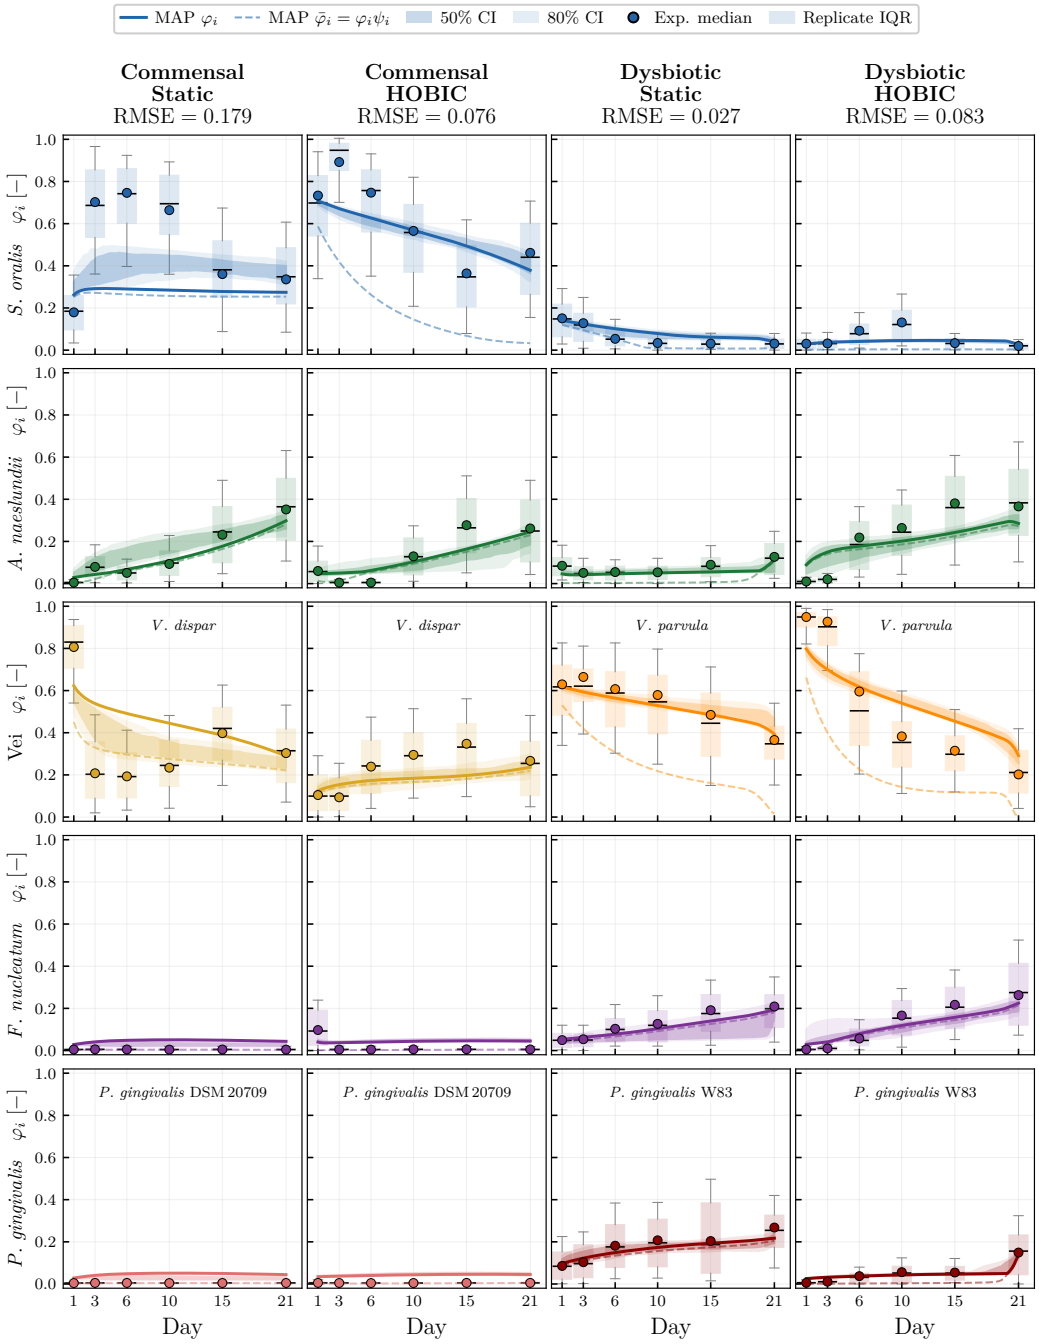


Figure 1: Species fractions over 21 days in the four conditions (columns) for the five taxa (rows). Lines: MAP of stage (iv), species volume fraction ϕ_i (solid) and living fraction $\bar{\phi}_i$ (dashed); bands: 50% and 80% posterior predictive intervals pooled over three seeds; dots and boxes: median and interquartile range of the replicates [8]. The RMSE of the MAP is given above each column. **[DS and DH from stage (ii) until stage (iv) is accepted.]**