Supplementary Material

Identifiability of species interactions in a living material: Bayesian calibration of an extended Hamilton principle multiphase model

K. Nishioka, F. Klempt, H. Geisler, R. Mukherjee, N. Heine, K. Doll-Nikutta, M. Stiesch, S. P. Szafrański, M. Soleimani, P. Junker

S1 Model derivation

State variables and kinematics

Following Klempt et al. (2026, 2024), the biofilm is described at a material point by two sets of internal variables: the volume fraction $\phi_i \in [0,1]$ occupied by species i ($i = 1, \dots, n$; here $n = 5$), and the viability fraction $\psi_i \in [0,1]$ representing the proportion of living cells within species i. The living volume fraction is then

$$\bar{\phi}_i := \phi_i \, \psi_i, \qquad i = 1, \dots, n. \tag{S1}$$

An additional void fraction ϕ_0 closes the system via the holonomic constraint

$$\sum_{l=0}^{n} \phi_l(t) - 1 = 0 \quad \forall \, t. \tag{S2}$$

Variational formulation

The evolution equations are derived from the principle of the minimum of the dissipation potential (Junker et al., 2014; Hackl, 1997), the local form of the extended Hamilton principle (Junker and Balzani, 2021; Junker et al., 2025):

$$\dot{\Psi} + \Delta + C \to \text{stat}_{(\dot{\boldsymbol{\xi}}, \gamma)}, \tag{S3}$$

where $\boldsymbol{\xi} = (\boldsymbol{\phi}, \boldsymbol{\psi})^T$ collects the volume fractions and viability fractions, γ is a Lagrange multiplier enforcing the volume constraint, C is the constraint function, and Δ the dissipation function. The free energy density encodes nutrient-driven growth and antibiotic-induced killing:

$$\Psi = -\tfrac{1}{2} \, c^* \, \bar{\boldsymbol{\phi}}^T \mathbf{A} \, \bar{\boldsymbol{\phi}} + \tfrac{1}{2} \, \alpha^* \, \boldsymbol{\psi}^T \mathbf{B} \, \boldsymbol{\psi}, \tag{S4}$$

where c^* is the nutrient concentration, α^* the antibiotic concentration, $\mathbf{A} \in \mathbb{R}^{n \times n}$ the symmetric growth coefficient matrix, and $\mathbf{B} = \text{diag}(b_1, \dots, b_n)$ the antibiotic sensitivity matrix.

The dissipation function is derived from the rate-dependent potential (Klempt et al., 2026):

$$\Delta = \tfrac{1}{2} \, \dot{\bar{\boldsymbol{\phi}}}^T \boldsymbol{\eta} \, \dot{\bar{\boldsymbol{\phi}}} + \tfrac{1}{2} \, \dot{\boldsymbol{\phi}}^T \boldsymbol{\eta} \, \dot{\boldsymbol{\phi}}, \tag{S5}$$

with the diagonal viscosity matrix $\boldsymbol{\eta} = \text{diag}(\eta_1, \dots, \eta_n)$. The first term couples $\dot{\phi}_i$ and $\dot{\psi}_i$ through $\dot{\bar{\phi}}_i = \dot{\phi}_i \psi_i + \phi_i \dot{\psi}_i$, giving rise to non-linear dissipative interactions; the second term regularises the system to full rank. Following Klempt et al. (2026), the viscosity parameters are set to $\eta_i = 1$ for all species, rendering them dimensionless scaling factors; the physical time scale is absorbed into the time-step size Δt.

The volume constraint (S2) enters through a Lagrange multiplier γ:

$$C = \gamma \Big(\sum_{l=0}^{n} \phi_l - 1 \Big). \tag{S6}$$

Stationarity conditions

Evaluating the stationarity conditions yields the coupled evolution equations (Klempt et al., 2026):

Volume fraction ($i = 1, \ldots, n$):

$$0 = -c^* \, \psi_i \left(\mathbf{A}\bar{\boldsymbol{\phi}}\right)_i + \eta_i(\dot{\phi}_i \, \psi_i^2 + \bar{\phi}_i \, \dot{\psi}_i + \dot{\phi}_i) + \gamma, \tag{S7}$$

Viability fraction ($i = 1, \ldots, n$):

$$0 = -c^* \, \phi_i \left(\mathbf{A}\bar{\boldsymbol{\phi}}\right)_i + \alpha^* \, b_i \, \psi_i + \eta_i(\dot{\psi}_i \, \phi_i^2 + \bar{\phi}_i \, \dot{\phi}_i) + \gamma, \tag{S8}$$

Void fraction:

$$0 = \gamma + \dot{\phi}_0, \tag{S9}$$

Volume constraint:

$$0 = \sum_{l=0}^{n} \phi_l - 1, \tag{S10}$$

In the numerical implementation, a logarithmic barrier penalty Ψ_{bar} with $K_p = \eta_i \times 10^{-4}$ is added to enforce the box constraints $\phi_i, \psi_i \in (0, 1)$ (Klempt et al., 2026); its contribution is omitted from Eqs. (S7)–(S9) for clarity. The system (S7)–(S10) constitutes $2n + 2$ coupled non-linear equations for the unknowns $(\phi_1, \ldots, \phi_n, \phi_0, \psi_1, \ldots, \psi_n, \gamma)$, solved at each time step via Newton's method with an implicit Euler discretisation ($\Delta t = 10^{-4}$, $N_{\text{step}} = 2{,}500$ steps, $t_{\max} = 0.25$; day d is mapped to $t = 0.2375\, d/21$, so day 21 lies at 95% of the integration window).

Only the symmetric part of $\mathbf{A}$ enters the quadratic free energy, so $\mathbf{A}$ is taken symmetric, $a_{ij} = a_{ji}$; the variational derivation ensures thermodynamic consistency (Junker and Balzani, 2021). Since c^* enters the evolution equations only through the product $c^* a_{ij}$, its absolute value is non-identifiable; we fix $c^* = 25$ so that the inferred a_{ij} remain $\mathcal{O}(1)$. The Heine et al. experiments do not involve antibiotic agents, so we set $\alpha^* = 0$, which eliminates the decay term $\alpha^* b_i \psi_i$ from Eq. (S8); the antibiotic sensitivity matrix $\mathbf{B}$ is therefore inactive and excluded from estimation. The symmetric interaction matrix $\mathbf{A}$ has 15 independent entries (the main text); these form the 15-dimensional parameter vector $\boldsymbol{\theta}$.

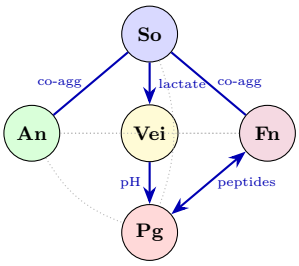
S2 Reported interactions among the five taxa

Figure S1 and Table S1 summarise interactions reported in the literature; they are used only to interpret the inferred coefficients, never as constraints.

S3 Likelihood weights and algorithm

Weight	Value	Applies to
λ_{Pg}	5	Pg, all time points
λ_{late}	3	Days 15 and 21, all species
λ_{rare}	0.1	species with mean fraction $< 5\%$
λ_{ψ}	3.0 (DH), 2.0 (DS), 1.5 (CS, CH)	viability channel, Phase 2

Table S2: Likelihood weights.



Pair	Mechanism	Reported flow
So–An	Co-aggregation	–
So–Vei	Lactate exchange	So to Vei
So–Fn	Co-aggregation (Fn as bridging organism)	–
Vei–Pg	pH elevation	Vei to Pg
Fn–Pg	Peptide supply	mutual

Table S1: Previously reported or biologically supported interactions (Kolenbrander et al., 2010; Heine et al., 2025) used to interpret the inferred coefficients. The model coefficients a_{ij} are symmetric; the last column refers to the reported metabolite flow (–: physical co-aggregation, no metabolite flow), not to the model. All 15 parameters are estimated; none is fixed a priori.

Figure S1: Interactions considered in interpreting the inferred coefficients, compiled from the literature (Kolenbrander et al., 2010; Heine et al., 2025). Solid: previously reported or biologically supported interaction (labels: reported mechanism; arrows: reported metabolite flow; lines without arrows: co-aggregation; neither is a property of the model, whose matrix $\mathbf{A}$ is symmetric). Dotted grey: no interaction considered in the interpretation; this does not mean that none exists. The network is not used in the inference: all 15 entries of $\mathbf{A}$ are estimated with the same uniform prior box.

Metric	CPU (12 cores)	GPU (RTX 4090)	Speedup
Matched benchmark (DH, $N_p = 500$, $K = 10$, 9 stages):			
Per stage	2 400 s	14 s	170×
Total	21 600 s	147 s	**147**×

Table S3: Computational cost comparison. The matched benchmark (500 particles, identical settings) isolates the GPU speed-up at about 150× (170× per tempering step, 147× end to end). The cost of every inference stage of the paper is listed in Table S10.

Algorithm S1 Staged GPU-accelerated TMCMC (one seed; every stage is run with three seeds and checked as in the main text)

Require: Composition data $\mathbf{y}_{\text{obs}}$, viability data, N_p particles, GPU device
Ensure: MAP estimate $\hat{\boldsymbol{\theta}}^{\text{MAP}}$, posterior samples, log-evidence $\ln \hat{Z}$

1: ▷ **Stage (i): composition likelihood, initial box**
2: Draw $\{\boldsymbol{\theta}_j\}_{j=1}^{N_p}$ from the stage prior; set $\beta_0 \leftarrow 0$
3: **while** $\beta_m < 1$ **do**
4: Find β_{m+1} s.t. $\mathrm{CoV}[\{w_j\}] = \delta_{\text{target}}$ via bisection
5: Compute weights $w_j \leftarrow p(\mathbf{y}_{\text{obs}} \mid \boldsymbol{\theta}_j)^{\beta_{m+1}-\beta_m}$
6: Resample N_p particles $\propto w_j$ (systematic)
7: Estimate weighted covariance $\hat{\boldsymbol{\Sigma}}^{(m)}$ (weighted particle covariance)
8: **for** $k = 1$ **to** K **do**
9: Propose $\boldsymbol{\theta}'_j$ for all j (random-walk or differential-evolution move)
10: Evaluate all N_p likelihoods in parallel via `jax.vmap` on GPU
11: Accept/reject via MH ; reject proposals outside the box
12: **end for**
13: **end while**
14: Record MAP and posterior SD per parameter
15: ▷ **Stage (ii): composition likelihood, box mean ± 4 SD**
16: Repeat lines 2–13 on the narrowed box, starting from the previous particles
17: ▷ **Stages (iii)–(iv): add the viability channel; initial box, then mean ± 4 SD**
18: Augment likelihood with the weighted viability channel (Table S2)
19: Repeat lines 2–13 for each stage
20: **return** $\hat{\boldsymbol{\theta}}^{\text{MAP}}$, posterior samples, $\ln \hat{Z}$

S4 Prior boxes

Table S4 lists the uniform prior box of every interaction entry for each condition and stage. The initial box was set empirically to contain earlier posterior estimates and encodes no sign constraint. An entry was widened only where more than 20% of its posterior mass lay in the outer 5% of the range on one side, or where this fraction increased by more than 0.10 between stages with the same likelihood; the stage was then repeated. The last column gives $\ln V = \sum_k \ln(u_k - l_k)$, the log-volume of the box, which enters $\ln \hat{Z}$ and is the reason why $\ln \hat{Z}$ is not compared across conditions.

S5 Runs and convergence

Table S5 lists, for every condition and stage, the number of particles and mutation steps, the number of tempering stages, the acceptance rate, the moves per parameter, and the spread over the three seeds of the maximum log-likelihood and of the posterior medians (criteria (a)–(e) of the main text).

S6 Identifiability of the interaction entries

Table S6 lists, for the final stage (iv) of every condition, the median and the 90% interval of each entry, pooled over the three seeds. An entry whose posterior standard deviation is at least 0.8 times that of a uniform distribution on its box is shown in grey and excluded from the seed-agreement criterion: the data do not constrain its magnitude beyond the prior box. Among these, entries whose 90% interval excludes zero are marked s.o. (sign only), the others n.i. (not

Table S4: Prior boxes per condition and stage for the 15 free entries. (i) initial box; (ii) stage-(i) posterior mean ± 4 SD; (iii) back on the box of stage (i), with the bounds widened (bold) where more than 20% of the posterior mass lay in the outer 5% on one side; (iv) stage-(iii) posterior mean ± 4 SD, clipped to the stage-(iii) box (seed with the highest log-likelihood). $\ln V = \sum_k \ln(u_k - l_k)$ over the free entries. **[DS, DH: stage (iv) pending.]**

CS

Entry	(i)	(ii)	(iii)	(iv)
a_{11}	[0, 2]	[0, 2]	**[-5, 2]**	[-5, 1.28]
a_{12}	[-0.5, 2]	[-0.18, 2]	[-0.5, 2]	[-0.5, 2]
a_{22}	[-1, 3]	[-1, 3]	[-1, 3]	[-1, 3]
a_{33}	[-1, 1.5]	[-1, 0.95]	**[-15, 20]**	[-15, 6.12]
a_{34}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{44}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{13}	[-1, 3]	[1.17, 3]	**[-1, 20]**	[-1, 13.99]
a_{14}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{23}	[-1, 3]	[-0.71, 3]	[-1, 3]	[-1, 3]
a_{24}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{55}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{15}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{25}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{35}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{45}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
$\ln V$	12.9	11.7	18.5	17.5

CH

Entry	(i)	(ii)	(iii)	(iv)
a_{11}	[-1, 1.5]	[-1, 1.5]	**[-5, 1.5]**	[-4.18, 1.5]
a_{12}	[-1, 2.5]	[-0.84, 2.5]	**[-5, 2.5]**	[-5, 2.5]
a_{22}	[-1, 2.5]	[-1, 2.5]	[-1, 2.5]	[-1, 2.5]
a_{33}	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]	[-1, 1.03]
a_{34}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{44}	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]
a_{13}	[-1, 3]	[-0.4, 3]	[-1, 3]	[-1, 2.36]
a_{14}	[-1, 1]	[-1, 1]	**[-5, 1]**	[-5, 1]
a_{23}	[-1, 3]	[-0.85, 3]	[-1, 3]	[-1, 3]
a_{24}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{55}	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]
a_{15}	[-1, 1]	[-1, 1]	**[-5, 1]**	[-5, 1]
a_{25}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{35}	[-1, 1]	[-1, 1]	[-1, 1]	[-1, 1]
a_{45}	[-1, 2]	[-1, 2]	[-1, 2]	[-1, 2]
$\ln V$	14.2	14.0	18.1	17.6

DS *(stage (iii) is final, provisional)*

Entry	(i)	(ii)	(iii)	(iv)
a_{11}	[-1, 2]	[-1, 2]	[-1, 2]	–
a_{12}	[-1, 1]	[-1, 1]	[-1, 1]	–
a_{22}	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]	–
a_{33}	[-15, 20]	[-15, 15.35]	[-15, 20]	–
a_{34}	[-15, 20]	[-10.47, 11.48]	[-15, 20]	–
a_{44}	[-1, 2.5]	[-1, 2.5]	[-1, 2.5]	–
a_{13}	[-1, 1.5]	[-1, 1.5]	[-1, 1.5]	–
a_{14}	[-1, 1]	[-1, 1]	[-1, 1]	–
a_{23}	[-1, 3]	[-1, 3]	**[-5, 3]**	–
a_{24}	[-0.5, 3]	[-0.5, 3]	[-0.5, 3]	–
a_{55}	[-1, 4]	[-1, 2.49]	**[-5, 4]**	–
a_{15}	[-0.5, 2.5]	[-0.5, 2.5]	**[-5, 2.5]**	–
a_{25}	[-0.5, 2.5]	[-0.5, 2.5]	[-0.5, 2.5]	–
a_{35}	[-3, 3]	[-2.04, 3]	[-3, 3]	–
a_{45}	[-15, 20]	[-2.92, 7.26]	[-15, 20]	–
$\ln V$	24.5	22.1	26.7	–

DH *(stage (iii) is final, provisional)*

Entry	(i)	(ii)	(iii)	(iv)
a_{11}	[0, 5]	[0, 5]	**[-5, 5]**	–
a_{12}	[0.5, 3]	[0.5, 3]	**[-5, 3]**	–
a_{22}	[0, 5]	[0, 2.02]	**[-5, 5]**	–
a_{33}	[-1, 3]	[-0.66, 1.73]	**[-15, 20]**	–
a_{34}	[-1, 1]	[-0.28, 1]	**[-15, 20]**	–
a_{44}	[0, 5]	[0, 3.3]	**[-10, 5]**	–
a_{13}	[-3, 3]	[-3, 3]	[-3, 3]	–
a_{14}	[-1.5, 3]	[-1.5, 3]	**[-8, 3]**	–
a_{23}	[0, 3]	[0, 3]	**[-15, 20]**	–
a_{24}	[-0.5, 3]	[-0.5, 3]	**[-15, 20]**	–
a_{55}	[-1.5, 5]	[-1.5, 5]	[-1.5, 5]	–
a_{15}	[-0.5, 2.5]	[-0.5, 2.5]	[-0.5, 2.5]	–
a_{25}	[-0.5, 2.5]	[-0.5, 2.5]	[-0.5, 2.5]	–
a_{35}	[-0.5, 5]	[-0.5, 2.49]	**[-15, 20]**	–
a_{45}	[-0.5, 6]	[-0.5, 6]	**[-5, 20]**	–
$\ln V$	21.1	18.2	38.6	–

identified).

Curvature of the likelihood. As a check that does not use the posterior samples, we computed the Hessian $\mathbf{H}$ of $-\ln p(\mathbf{y}_{\text{obs}} \mid \boldsymbol{\theta})$ at the MAP over the 15 free entries, scaled each entry by the standard deviation of a uniform distribution on its box, and counted the eigen-directions with scaled curvature of at least $1/0.8^2$, i.e. directions in which the likelihood is narrower than the identification threshold. Their number is 4, 4, 8 and 5 in CS, CH, DS and DH, against 4, 6, 8 and 9 entries identified from the posterior **[prov.]**. The counts agree in CS and DS; in CH and DH the posterior identifies more entries than the local curvature, because several entries are constrained jointly along correlated directions and by the non-quadratic shape of the likelihood, which a single Hessian at the MAP does not capture. The MAP is the best particle and not an exact optimum (4–5 small negative eigenvalues), so entry-wise Laplace intervals are not reported.

Table S5: Runs of the final pipeline: particles N_p, mutation steps K, tempering stages, acceptance rate, moves per free parameter, maximum log-likelihood (best seed; spread over three seeds in parentheses) and wall-clock hours per seed.

Cond.	Stage	N_p	K	Steps	Accept.	Moves/dim	max log L (spread)	Hours
CS	(i)	1000	80	6	0.14–0.14	4.4–4.5	-5.86 (0.11)	2.9–3.0
CS	(ii)	2000	120	6	0.14–0.14	6.7–6.8	-5.67 (0.27)	4.7–5.1
CS	(iv)	5000	150	6–7	0.13–0.13	7.9–9.2	-15.81 (0.24)	7.7–9.0
CH	(i)	1000	80	5	0.14–0.14	3.8–3.9	-3.19 (0.10)	2.5–2.5
CH	(ii)	2000	150	4	0.15–0.15	5.9–6.0	-3.15 (0.08)	3.9–4.2
CH	(iv)	5000	150	7	0.13–0.13	8.9–9.0	-8.65 (0.11)	9.7–9.8
DS	(i)	1000	80	10	0.11–0.12	5.9–6.2	-0.91 (0.72)	5.2–5.4
DS	(ii)	3000	100	10	0.12–0.12	8.0–8.2	-0.58 (0.27)	5.7–6.4
DS	(iii)	18000	150	10	0.11–0.12	11.4–11.6	-6.45 (0.53)	15.8–17.6
DH	(i)	1000	80	10	0.12–0.13	6.6–6.7	-3.92 (0.51)	4.9–5.1
DH	(ii)	2000	80	8–9	0.12–0.13	5.4–6.0	-3.21 (0.31)	4.3–5.2
DH	(iii)	14000	120	12	0.10–0.10	9.9–10.1	-14.65 (0.16)	15.3–16.6

Table S6: Posterior median [5%, 95%] of the 15 entries at stage (iv), three seeds pooled. Grey: the posterior spreads over most of the box (standard deviation ≥ 0.8 times that of a uniform distribution on the box); n.i.: not identified, s.o.: sign only (the 90% interval excludes zero).

Entry	CS	CH	DS	DH
a_{11}	−3.58 [−4.9, −1.2]	−0.99 [−2.5, −0.5]	+0.02 [−0.9, +1.7] n.i.	−0.85 [−4.6, +4.3] n.i.
a_{12}	+0.77 [−0.3, +1.8] n.i.	−0.44 [−3.5, +0.6]	−0.07 [−0.9, +0.9] n.i.	−1.35 [−4.6, +2.3] n.i.
a_{22}	+0.68 [−0.7, +2.2]	−0.13 [−0.9, +1.5]	+0.14 [−0.9, +1.3] n.i.	−2.98 [−4.7, −0.7]
a_{33}	−6.76 [−13.4, −2.5]	−0.53 [−1.0, +0.2]	−0.63 [−1.5, −0.3]	−1.28 [−2.8, −0.7]
a_{34}	−0.29 [−0.9, +0.8] n.i.	−0.20 [−0.9, +0.8] n.i.	−2.25 [−12.3, +0.5]	−1.29 [−4.4, +3.7]
a_{44}	−0.01 [−0.9, +0.9] n.i.	+0.19 [−0.9, +1.4] n.i.	−0.12 [−0.9, +2.0] n.i.	−6.88 [−9.7, −2.4]
a_{13}	+5.74 [+2.8, +9.2]	+0.61 [+0.0, +1.4]	−0.16 [−0.9, +0.7]	−1.00 [−2.7, +1.1]
a_{14}	−0.40 [−0.9, +0.6] n.i.	−2.89 [−4.8, −0.2] s.o.	−0.10 [−0.9, +0.9] n.i.	−3.19 [−7.0, +1.3] n.i.
a_{23}	+0.50 [−0.8, +2.5] n.i.	+0.98 [−0.1, +2.0]	−1.84 [−4.5, −0.1]	+1.18 [−2.1, +2.9]
a_{24}	−0.22 [−0.9, +0.8] n.i.	−0.13 [−0.9, +0.8] n.i.	+1.16 [−0.3, +2.8] n.i.	+5.10 [+2.7, +7.1]
a_{55}	−0.04 [−0.9, +0.9] n.i.	+0.18 [−0.9, +1.4] n.i.	−2.09 [−4.7, −0.3]	+0.97 [−1.2, +3.8] n.i.
a_{15}	−0.42 [−0.9, +0.6] n.i.	−3.12 [−4.8, −0.3] s.o.	−1.45 [−4.3, +0.5]	+0.99 [−0.2, +2.3] n.i.
a_{25}	−0.23 [−0.9, +0.8] n.i.	−0.20 [−0.9, +0.8] n.i.	+1.18 [−0.3, +2.4] n.i.	+1.10 [−0.2, +2.3] n.i.
a_{35}	−0.32 [−0.9, +0.7] n.i.	−0.29 [−0.9, +0.8] n.i.	+1.04 [+0.3, +2.4]	−6.55 [−13.4, −1.7]
a_{45}	−0.05 [−0.9, +0.9] n.i.	+0.45 [−0.9, +1.8] n.i.	+2.81 [−1.3, +15.1]	+1.21 [−1.9, +7.5]

S7 MAP and posterior bulk by seed

Figure S2 overlays the MAP of each seed on the posterior predictive band pooled over the three seeds. In all conditions the three MAPs nearly coincide. In CS the MAP lies apart from the bulk of the posterior for So and Vei: the pointwise posterior median has a composition error of 0.125, the MAP of 0.179; in the other conditions the two differ by less than 0.02**[prov.]**.

S8 Bayes factor for the *Fn–Pg* interaction

Stage (i) was repeated with $a_{45} = 0$ for all four conditions (three seeds each). Table S7 reports $\ln \hat{Z}$ of both models, $\ln B = \ln \hat{Z}_{\text{full}} - \ln \hat{Z}_{a_{45}=0}$, and the width of the a_{45} box, which sets the Occam penalty.

Table S7: Log-evidence of the full and the reduced ($a_{45} = 0$) model at stage (i), mean ± half-range over three seeds, and $\ln B = \ln \hat{Z}_{\text{full}} - \ln \hat{Z}_{a_{45}=0}$. The last two columns give the a_{45} box and its log-width, which enters $\ln \hat{Z}_{\text{full}}$ as an Occam penalty.

Cond.	$\ln \hat{Z}_{\text{full}}$	$\ln \hat{Z}_{a_{45}=0}$	$\ln B$	a_{45} box	ln(width)
CS	-15.16 ± 0.08	-15.26 ± 0.13	$+0.10$	[-1, 1]	0.69
CH	-9.62 ± 0.12	-9.66 ± 0.00	$+0.05$	[-1, 2]	1.10
DS	-21.48 ± 0.70	-18.81 ± 0.51	-2.67	[-15, 20]	3.56
DH	-21.36 ± 0.38	-23.30 ± 0.34	$+1.94$	[-0.5, 6]	1.87

Dysbiotic Static. In DS, a_{45} and a_{33} trade off at stage (i). With $a_{45} = 0$ the posterior moves almost entirely to the mode $a_{33} \approx -10$ (weight 0.83–0.87 across seeds), whereas the full model splits its mass between $a_{33} \approx +1$ and $a_{33} < -5$, with a weight of the latter that differs strongly between seeds (0.14–0.66); this stage-(i) group does not pass criterion (e) and is used here only for the Bayes factor. In the mode $a_{33} \geq -5$ the coefficient a_{45} is larger than in the mode $a_{33} < -5$ (Table S8). At the final stage the mode $a_{33} < -5$ is empty in all three seeds, and a_{45} has median +2.8 with a 90% interval reaching zero ($[-1.3, +15.1]$)**[prov.]**.

Table S8: DS, stage (i), full model: a_{45} by a_{33} mode (three seeds pooled, 3000 particles).

Mode	Particles	a_{45} median [5%, 95%]	max log-likelihood
$a_{33} \geq -5$	1761	+2.4 [+0.6, +4.7]	−0.91
$a_{33} < -5$	1239	+0.8 [−0.3, +1.6]	−1.49

S9 Effect of the likelihood weights

Stage (i) for DH was repeated with $\lambda_{\text{Pg}} = \lambda_{\text{late}} = 1$ (1000 particles, 100 mutation steps, three seeds; all convergence criteria met). Table S9 compares it with the weighted run.

Table S9: DH, stage (i): weighted ($\lambda_{\text{Pg}} = 5$, $\lambda_{\text{late}} = 3$) and unweighted likelihood. Ranges over three seeds; the measured Pg ratio of day 21 to day 15 is 2.74.

	a_{45} median [5%, 95%]	Pg D21/D15	RMSE
Weighted	+2.9 to +3.2 [+0.5 to +1.1, +5.2 to +5.5]	2.4–2.9	0.058–0.065
Unweighted	+3.0 [−0.06 to +0.17, +5.6]	1.7–3.1	0.059–0.068

References

Hackl, K., 1997. Generalized standard media and variational principles in classical and finite elastoplasticity. Journal of the Mechanics and Physics of Solids 45, 667–688. doi:10.1016/S0022-5096(96)00110-X.

Heine, N., Bittroff, K., Szafrański, S.P., Duitscher, M., Behrens, W., Vollmer, C., Mikolai, C., Kommerein, N., Debener, N., Frings, K., Heisterkamp, A., Scheper, T., Torres-Mapa, M.L., Bahnemann, J., Stiesch, M., Doll-Nikutta, K., 2025. Influence of species composition and cultivation condition on peri-implant biofilm dysbiosis in vitro. Frontiers in Oral Health 6, 1649419. doi:10.3389/froh.2025.1649419.

Junker, P., Balzani, D., 2021. An extended Hamilton principle as unifying theory for coupled problems and dissipative microstructure evolution. Continuum Mechanics and Thermodynamics 33, 1931–1956. doi:10.1007/s00161-021-01017-z.

Junker, P., Bode, T., Hackl, K., 2025. On a holistic variational formulation for material modeling including dissipative evolution. Journal of the Mechanics and Physics of Solids 200, 106133. doi:10.1016/j.jmps.2025.106133.

Junker, P., Makowski, J., Hackl, K., 2014. The principle of the minimum of the dissipation potential for non-isothermal processes. Continuum Mechanics and Thermodynamics 26, 259–268. doi:10.1007/s00161-013-0299-4.

Klempt, F., Geisler, H., Soleimani, M., Junker, P., 2026. A continuum multi-species bacterial growth model with a novel interaction scheme. Archive of Applied Mechanics 96, 164. doi:10.1007/s00419-026-03160-y.

Klempt, F., Soleimani, M., Wriggers, P., Junker, P., 2024. A Hamilton principle-based model for diffusion-driven biofilm growth. Biomechanics and Modeling in Mechanobiology 23, 2091–2113. doi:10.1007/s10237-024-01883-x.

Kolenbrander, P.E., Palmer Jr., R.J., Periasamy, S., Jakubovics, N.S., 2010. Oral multispecies biofilm development and the key role of cell–cell distance. Nature Reviews Microbiology 8, 471–480.

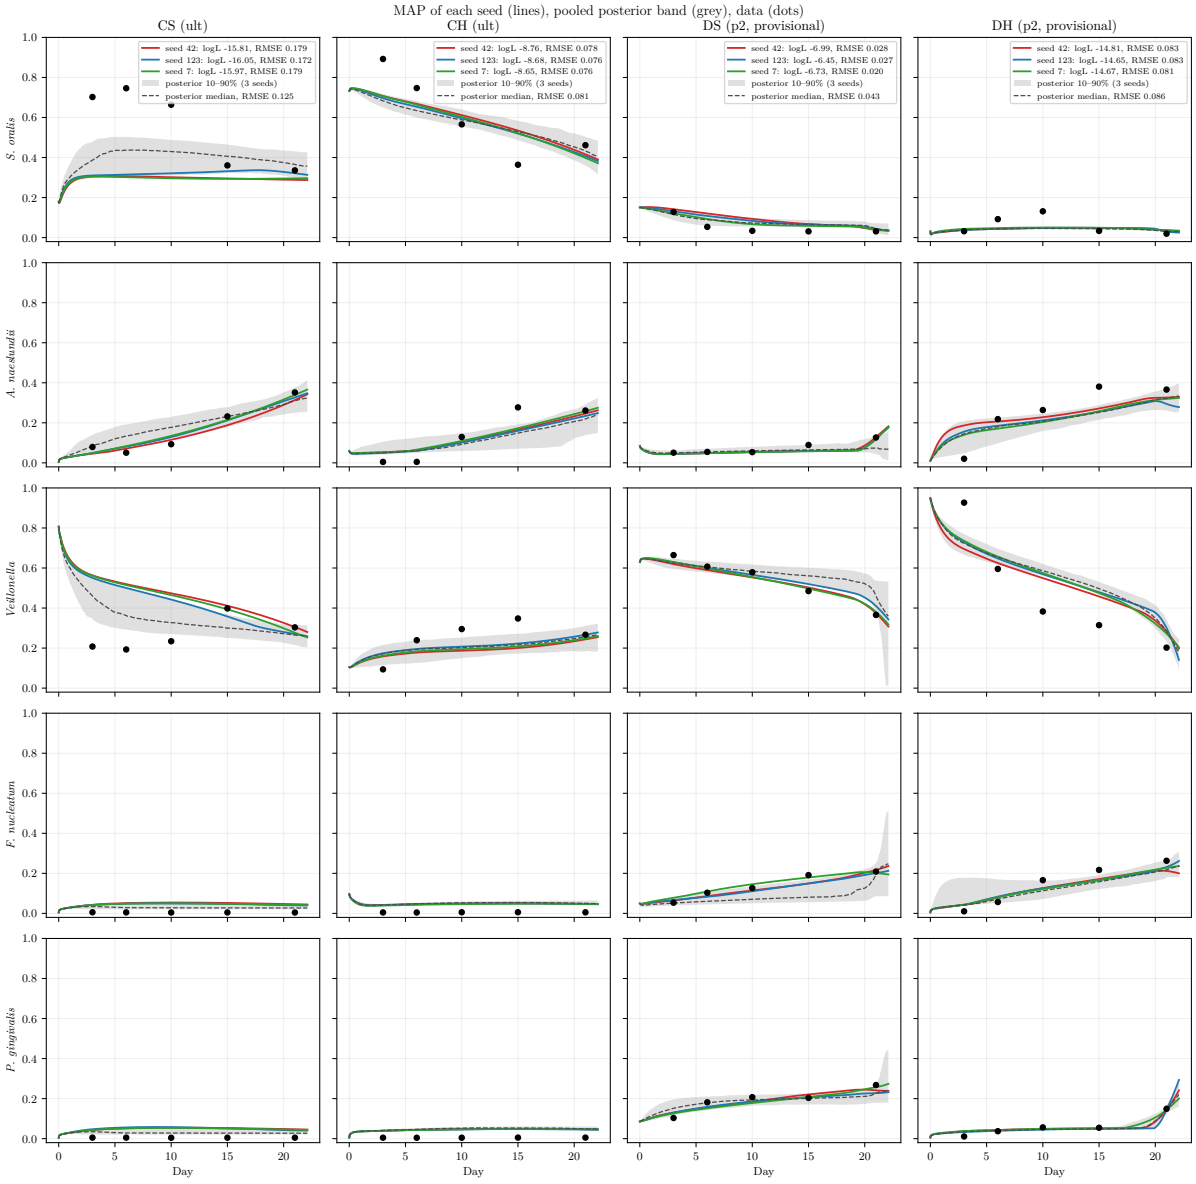


Figure S2: Species fractions: MAP of each seed (coloured lines), pooled posterior 10–90% band (grey) and its pointwise median (dashed), and data (dots). **[DS, DH from stage (iii).]**

Cond.	Stage	N_p	n_{mut}	K	Wall time / seed	N_{lik}
CS	(i)	1000	80	6	2.9 h	4.8×10^5
	(ii)	2000	120	6	4.7 h	1.4×10^6
	(iii)	2000	160	6	6.7 h	1.9×10^6
	(iv)	5000	150	7	8.9 h	5.3×10^6
CH	(i)	1000	80	5	2.5 h	4.0×10^5
	(ii)	2000	150	4	4.0 h	1.2×10^6
	(iii)	2000	150	7	7.4 h	2.1×10^6
	(iv)	5000	150	7	9.7 h	5.3×10^6
DS	(i)	1000	80	10	5.3 h	8.0×10^5
	(ii)	3000	100	10	5.7 h	3.0×10^6
	(iii)*	18000	150	10	17.6 h	2.7×10^7
DH	(i)	1000	80	10	5.0 h	8.0×10^5
	(ii)	2000	80	9	4.6 h	1.4×10^6
	(iii)*	14000	120	12	16.6 h	2.0×10^7

Table S10: Cost of each inference stage (median over the three seeds): particle number N_p, mutation steps per stage n_{mut}, number of tempering stages K, wall time per seed on one GPU, and number of likelihood evaluations $N_{\mathrm{lik}} = N_p(1 + \sum n_{\mathrm{mut}})$. *DS, DH: stage (iii), until stage (iv) is accepted **[prov.]**. Generated by `tools/make_cost_table.py`.